

# **CSE 575**

# **Statistical Machine Learning**

Lecture 9  
YooJung Choi  
Fall 2022

# Recap: Generative vs discriminative

$p(\mathbf{x}, \mathcal{C}_k)$  given by  
(Linear) Gaussian discriminant analysis  
Naïve Bayes classifier

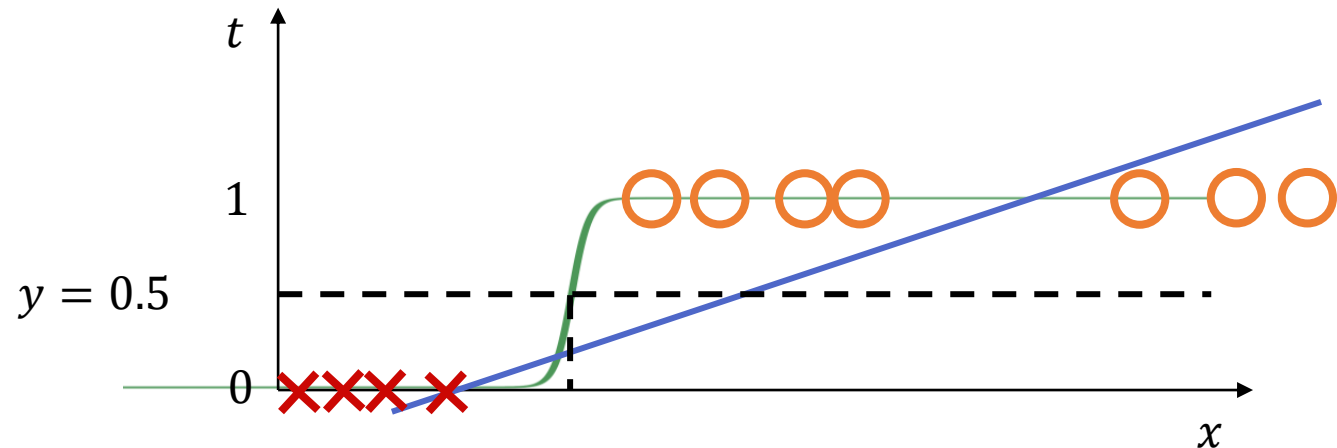
$p(\mathcal{C}_k|\mathbf{x})$  of the form  
 $\sigma(\mathbf{w}^T \mathbf{x} + w_0)$  (binary class), or  
 $s(\mathbf{w}^T \mathbf{x} + w_0)$  (multi-class)

- Can be used for various other tasks
- Performs very well *if* the modeling assumptions hold
- Closed-form solutions
- Tend to have more parameters

- Makes weaker assumptions—better performance in general
- Tend to have fewer parameters
- Limited to classification
- No closed-form solution

# Logistic regression

- Model  $p(t = 1|\mathbf{x})$  via  $y(\mathbf{x}) = \frac{1}{1+\exp\{-\mathbf{w}^T \mathbf{x}\}} = \sigma(\mathbf{w}^T \mathbf{x})$
- Again, assume  $x_0 = 1, t \in \{0,1\}$
- Recall: linear regression failed on this example



# Logistic regression

- Note:  $p(t = 1|\mathbf{x}) = y(\mathbf{x})$ , i.e.  $t \mid \mathbf{x} \sim \text{Bernoulli}(y(\mathbf{x}))$ .
- Given  $N$  data points  $\{(\mathbf{x}_n, t_n)\}$ , the likelihood function is:

$$p(\mathbf{t}|\mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n} \quad \text{where } y_n = y(\mathbf{x}_n) = \sigma(\mathbf{w}^T \mathbf{x}_n)$$

- Maximize log-likelihood, or equivalently, minimize the negative log-likelihood as the error function:

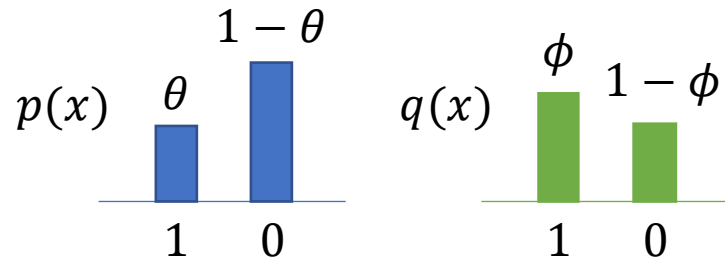
$$E(\mathbf{w}) = -\ln p(\mathbf{t}|\mathbf{w}) = -\sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\}$$

# Logistic regression

- From last slide: *Cross-entropy loss*

$$E(\mathbf{w}) = -\ln p(\mathbf{t}|\mathbf{w}) = -\sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\}$$

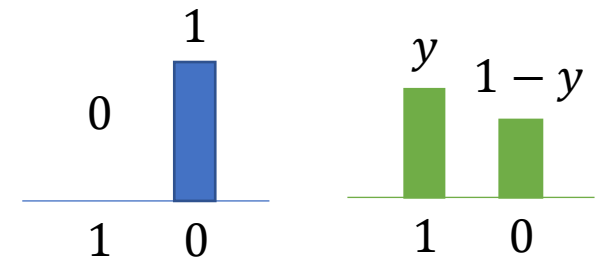
- Cross entropy of two distributions  $p$  and  $q$  over Boolean variables:



$$H(p, q) = -\theta \log \phi - (1 - \theta) \log(1 - \phi)$$

Quantifies how “close” the distributions are  
Minimized when  $p = q$

- $E(\mathbf{w})$ : Minimize the cross entropy between the ground-truth distribution and the distribution given by logistic regression



# Logistic regression

**Exercise:** derive  $\nabla E(\mathbf{w})$ .

Hint: first derive the derivative of logistic function  $\frac{d\sigma(a)}{da} = \sigma(a) \cdot (1 - \sigma(a))$

- Differentiate w.r.t.  $\mathbf{w}$ :

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \mathbf{x}_n$$

- Looks identical to the gradient of sum-of-squares error (linear regression)?
- $y_n = \sigma(\mathbf{w}^T \mathbf{x}_n)$  is *no longer linear*
- No known closed form solution
- Iteratively optimize via (stochastic) gradient descent!

$$\mathbf{w}^{(\text{new})} \leftarrow \mathbf{w}^{(\text{old})} - \eta \sum_{n=1}^N (y_n - t_n) \mathbf{x}_n$$

$$\mathbf{w}^{(\text{new})} \leftarrow \mathbf{w}^{(\text{old})} - \eta (y_n - t_n) \mathbf{x}_n$$

# Newton's method (or Newton-Raphson method)

- Iteratively find the root of a function  $f$ : i.e. find  $x$  s.t.  $f(x) = 0$
- Can be used to *minimize*  $E(\mathbf{w})$  by finding the root of  $\nabla E(\mathbf{w})$
- Compared to gradient descent, tends to converge in much *fewer iterations*.
- Each iteration is *more expensive*

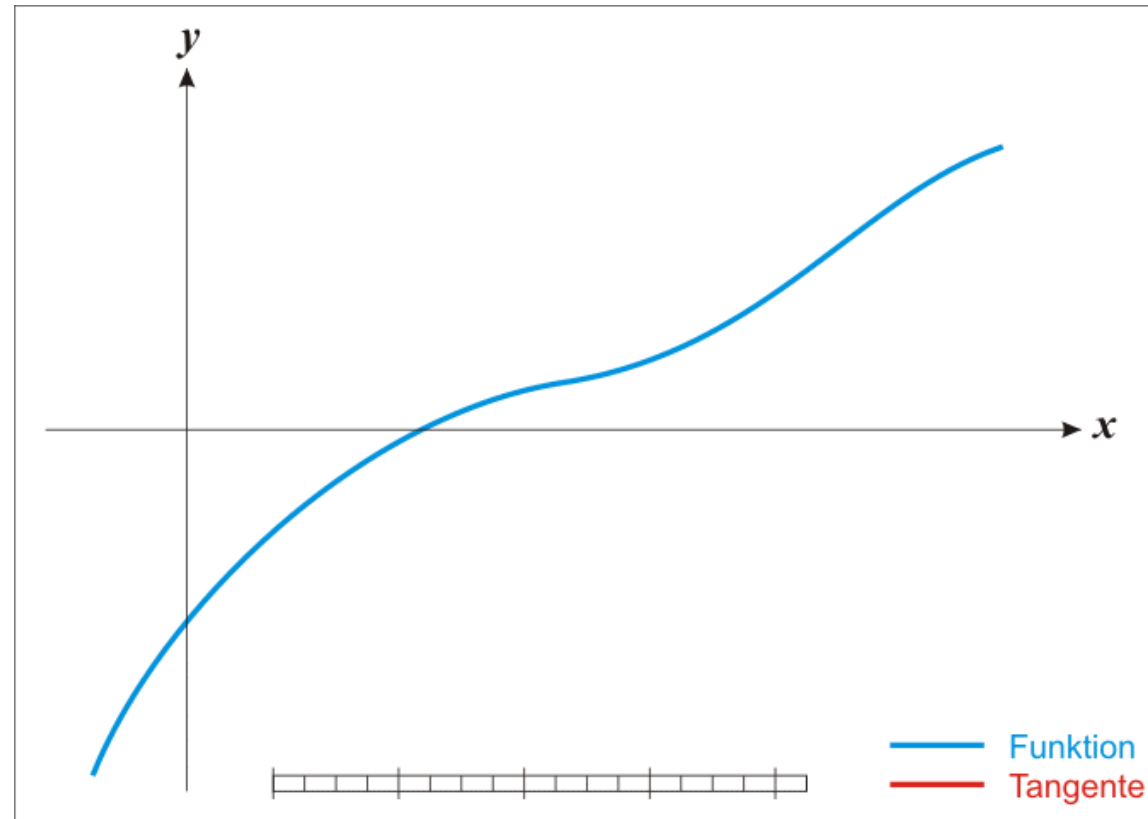
# Newton's method

- Iteratively: find the root of a linear approximation of  $f$



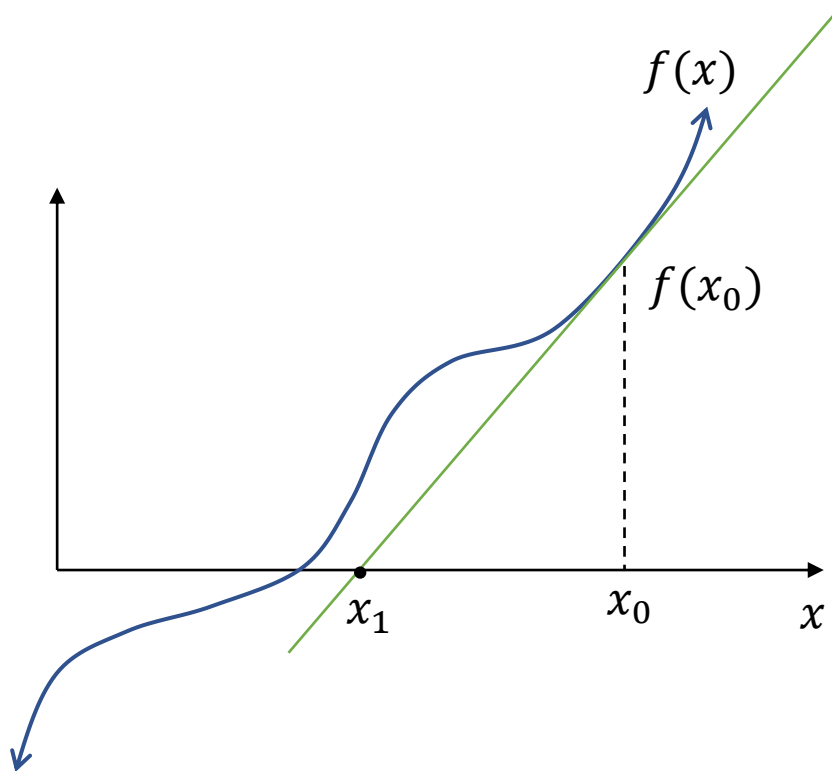
# Newton's method

- Iteratively: find the root of a linear approximation of  $f$



# Newton's method

- Iteratively: find the root of a linear approximation of  $f$



$$\text{"Slope"} = f'(x_0) = \frac{f(x_0)}{x_0 - x_1} \Rightarrow x_1 = x_0 - \frac{f(x_0)}{f'(x_0)}$$

To optimize  $f$  (i.e. find the root of  $f'$ ):

$$x_{i+1} = x_i - \frac{f'(x_i)}{f''(x_i)}$$

For a  $D$  dimensional input  $\mathbf{x}$ :

$$\mathbf{x}_{i+1} = \mathbf{x}_i - (\nabla^2 f(\mathbf{x}_i))^{-1} \nabla f(\mathbf{x}_i)$$

# Newton's method for logistic regression

- Applying Newton's method to minimize  $E(\mathbf{w})$ :

$$\mathbf{w}^{(\text{new})} \leftarrow \mathbf{w}^{(\text{old})} - \mathbf{H}^{-1} \nabla E(\mathbf{w})$$

where  $\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \mathbf{x}_n = \mathbf{X}^T (\mathbf{y} - \mathbf{t})$   $O(ND)$  time

$$\mathbf{H} = \nabla^2 E(\mathbf{w}) = \sum_{n=1}^N y_n (1 - y_n) \mathbf{x}_n \mathbf{x}_n^T = \mathbf{X}^T \mathbf{R} \mathbf{X} \quad O(ND^2) \text{ time}$$

$$R_{nn} = y_n (1 - y_n) \quad \text{Inverse: } O(D^3) \text{ time}$$

$$\mathbf{X} = \begin{bmatrix} - & \mathbf{x}_1^T & - \\ & \vdots & \\ - & \mathbf{x}_N^T & - \end{bmatrix}, \quad \mathbf{y} - \mathbf{t} = \begin{bmatrix} y(\mathbf{x}_1) - \mathbf{t}_1 \\ \vdots \\ y(\mathbf{x}_N) - \mathbf{t}_N \end{bmatrix}, \quad \mathbf{R} = \begin{bmatrix} y(\mathbf{x}_1)(1 - y(\mathbf{x}_1)) & & \\ & \ddots & \\ & & y(\mathbf{x}_N)(1 - y(\mathbf{x}_N)) \end{bmatrix}$$

# Softmax regression

- i.e. multiclass logistic regression

- For each class  $k$ :  $p(\mathcal{C}_k|\mathbf{x}) = y_k(\mathbf{x}) = s(\mathbf{w}_k^T \mathbf{x})$

$$s(a_k) = \frac{\exp(a_k)}{\sum_j \exp(a_j)}$$

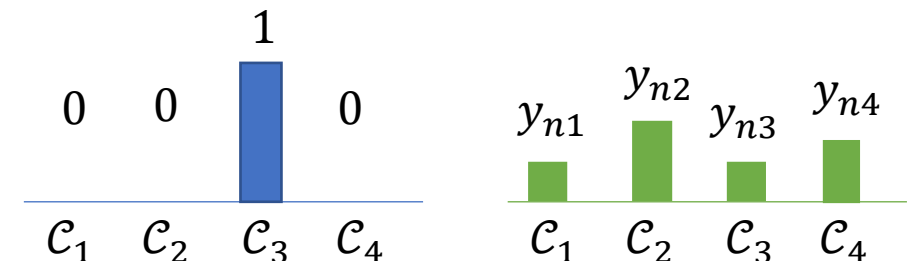
- Given  $N$  data points  $\{(\mathbf{x}_n, \mathbf{t}_n)\}$ , the likelihood function is:

$\mathbf{t}_n$ : 1-of-K (“one-hot”) vector

$$p(\mathbf{T}|\mathbf{w}_1, \dots, \mathbf{w}_K) = \prod_{n=1}^N \prod_{k=1}^K p(\mathcal{C}_k|\mathbf{x}_n)^{t_{nk}} = \prod_{n=1}^N \prod_{k=1}^K y_{nk}^{t_{nk}} \text{ where } y_{nk} = y_k(\mathbf{x}_n)$$

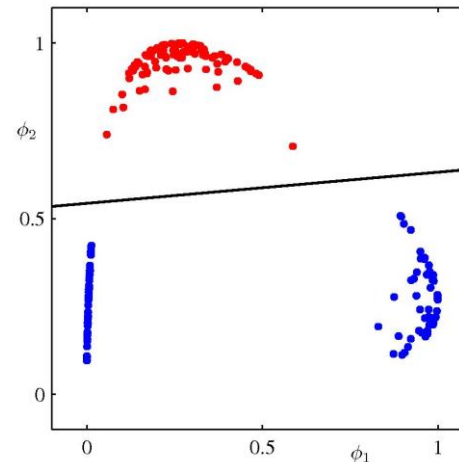
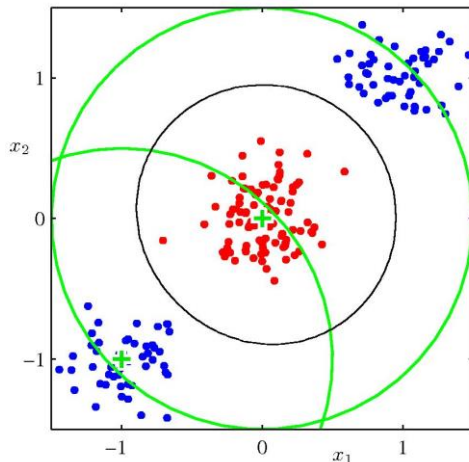
- Maximum likelihood = minimum log-likelihood = minimum cross-entropy loss:

$$E(\mathbf{w}_1, \dots, \mathbf{w}_K) = -\ln p(\mathbf{T}|\mathbf{w}_1, \dots, \mathbf{w}_K) = -\sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln y_{nk}$$



# Non-linear decision boundaries

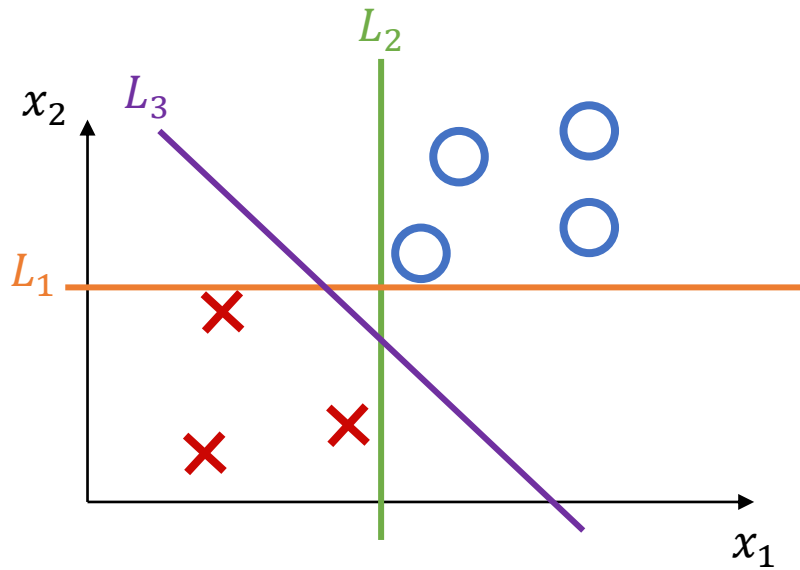
- Our discussion of linear regression  $\sigma(\mathbf{w}^T \mathbf{x})$  and softmax regression  $s(\mathbf{w}_k^T \mathbf{x})$  so far is limited to *linear* decision boundaries
- We can learn *non-linear* decision boundaries using fixed *non-linear* basis functions  $\Phi(\mathbf{x})$  (c.f. linear basis function models for regression)



*May require a very high dimensional feature space...*

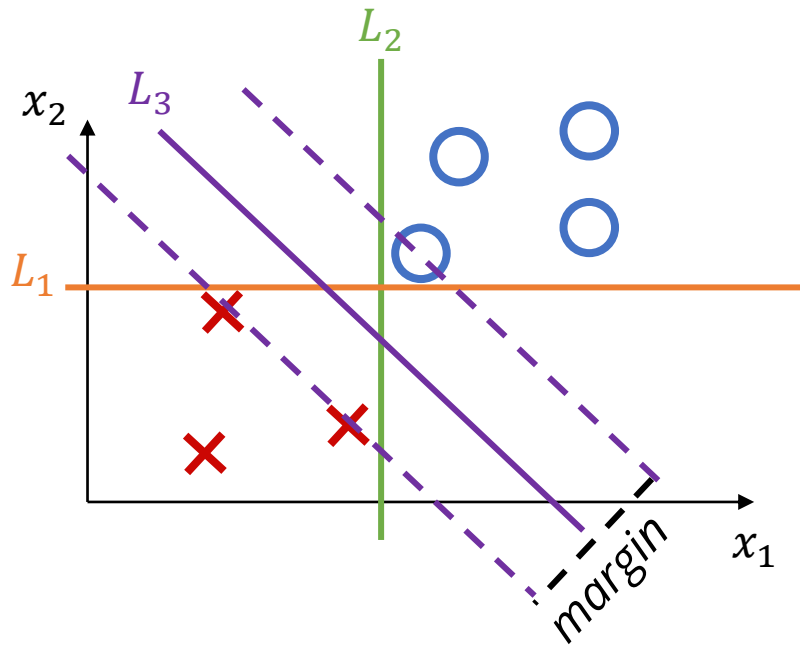
*Support vector machines handle this with the “kernel trick”*

# Optimal separating hyperplane



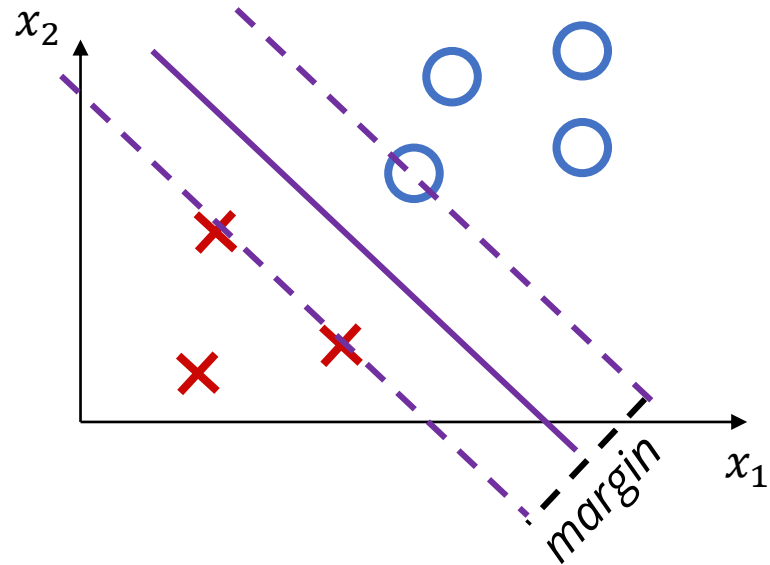
- Which decision boundary to choose?  $L_3$ 
  - $L_1$  and  $L_2$  just barely classify examples correctly. We would not expect them to generalize well
- How to design a learning algorithm that chooses  $L_3$  over  $L_1$  and  $L_2$ ?
  - E.g. you could end up with any of these boundaries using the perceptron algorithm

# Maximum margin classifier



- Idea: find the hyperplane with the most “cushion” between itself and the training examples
- Maximize the *margin*

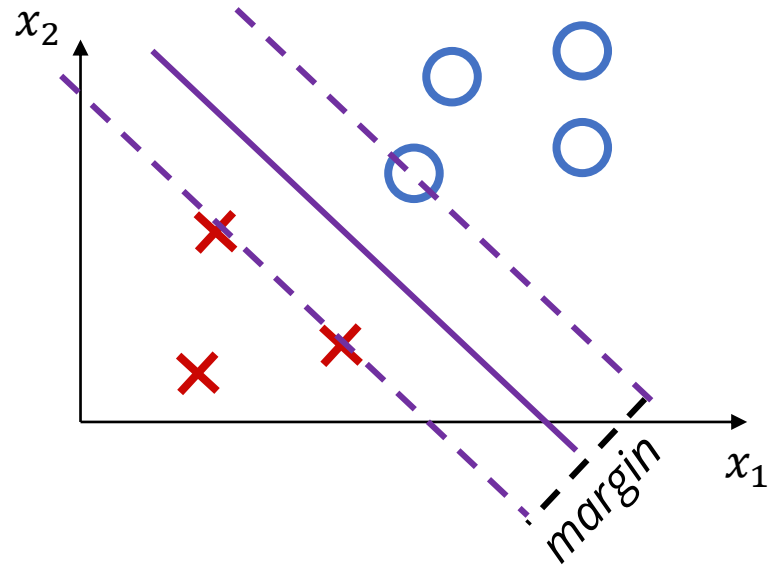
# Maximum margin classifier



- Assumptions & notations:
  - Binary class:  $t \in \{-1, +1\}$
  - Training data is linearly separable
  - Classification function:
$$\begin{cases} +1 & \text{if } \mathbf{w}^T \mathbf{x} + b \geq 0 \\ -1 & \text{if } \mathbf{w}^T \mathbf{x} + b < 0 \end{cases}$$
  - No longer assume a dummy feature  $x_0 = 1$



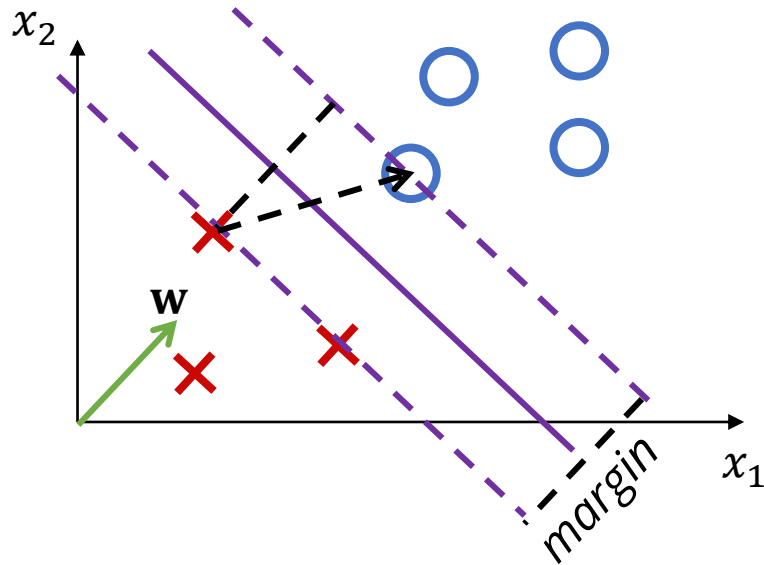
# Maximum margin classifier



- We want:
  - For every positive training example  $\mathbf{x}_+$ ,
$$\mathbf{w}^T \mathbf{x}_+ + b \geq 1$$
  - For every negative training example  $\mathbf{x}_-$ ,
$$\mathbf{w}^T \mathbf{x}_- + b \leq -1$$
- Equivalently, for every training example  $(\mathbf{x}_n, t_n)$ ,
$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1$$

# Maximum margin classifier

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \quad \forall n$$



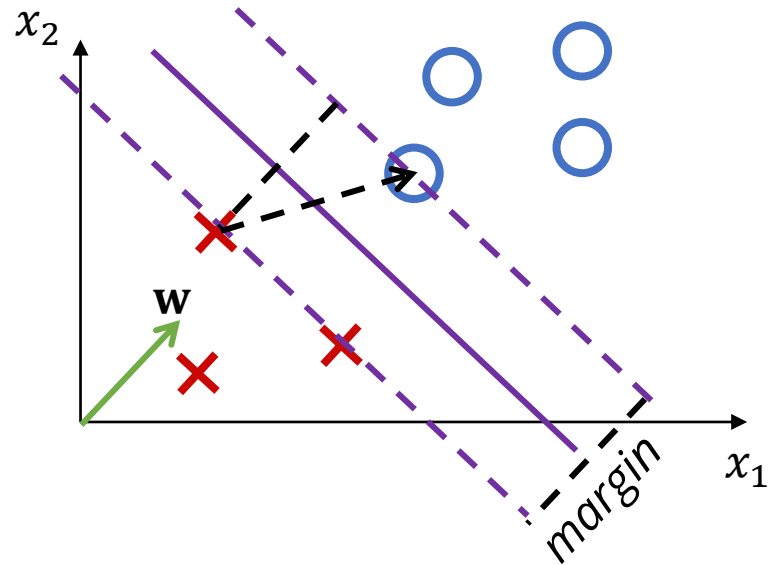
- If  $\mathbf{x}_+$  is the positive training example *closest to the hyperplane*, and  $\mathbf{x}_-$  the closest negative example,

$$\text{margin} = \frac{\mathbf{w}^T (\mathbf{x}_+ - \mathbf{x}_-)}{\|\mathbf{w}\|}$$

- Recall:  $\mathbf{w}^T \mathbf{x}_+ + b = 1$  and  $\mathbf{w}^T \mathbf{x}_- + b = -1$
- Then the margin is:  $\frac{2}{\|\mathbf{w}\|}$
- Maximum margin classifier:

$$\operatorname{argmax}_{\mathbf{w}, b} \frac{2}{\|\mathbf{w}\|} \quad \text{s.t. } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \quad \forall n$$

# Maximum margin classifier



$$\operatorname{argmax}_{\mathbf{w}, b} \frac{2}{\|\mathbf{w}\|} \quad \text{s.t. } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \quad \forall n$$

$$\operatorname{argmin}_{\mathbf{w}, b} \|\mathbf{w}\| \quad \text{s.t. } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \quad \forall n$$

$$\operatorname{argmin}_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{s.t. } t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 \quad \forall n$$

$$\operatorname{argmin}_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{s.t. } t_n(\mathbf{w}^T \mathbf{x}_n + b) - 1 \geq 0 \quad \forall n$$