

# CSE 594: Spatial Data Science & Engineering

Lecture 3  
Spatial SQL Part 1

# Database

- A large organized collection of data

## DBMS

- A software system to store, retrieve, and manipulate data
- Example – PostgreSQL, MySQL

## Relational Database

- A collection of structured data organized as a set of tables with rows and columns

# Why Storing Data in a Database?

Data is  
always  
synchronized

The diagram features three arrows pointing towards a central point. The top arrow is orange and points downwards. The bottom-left arrow is yellow and points upwards and to the right. The bottom-right arrow is grey and points upwards and to the left. The text is centered within each arrow.

Secure data from  
unauthorized  
access

Removes redundancy

# Why Relational Database?

- Easy to use
- Flexibility of making changes
- Concurrent collaboration among multiple users
- Compliance with ACID properties
- Reduces redundancy through normalization

# SQL – Structured Query Language

- A declarative programming language
  - Define only what to do
  - How to do is a black box
- Data Definition Language
  - Create, alter, delete tables and their attributes
- Data Manipulation Language
  - Retrieve, insert, delete, modify rows in the tables

# Schemas/Tables in SQL

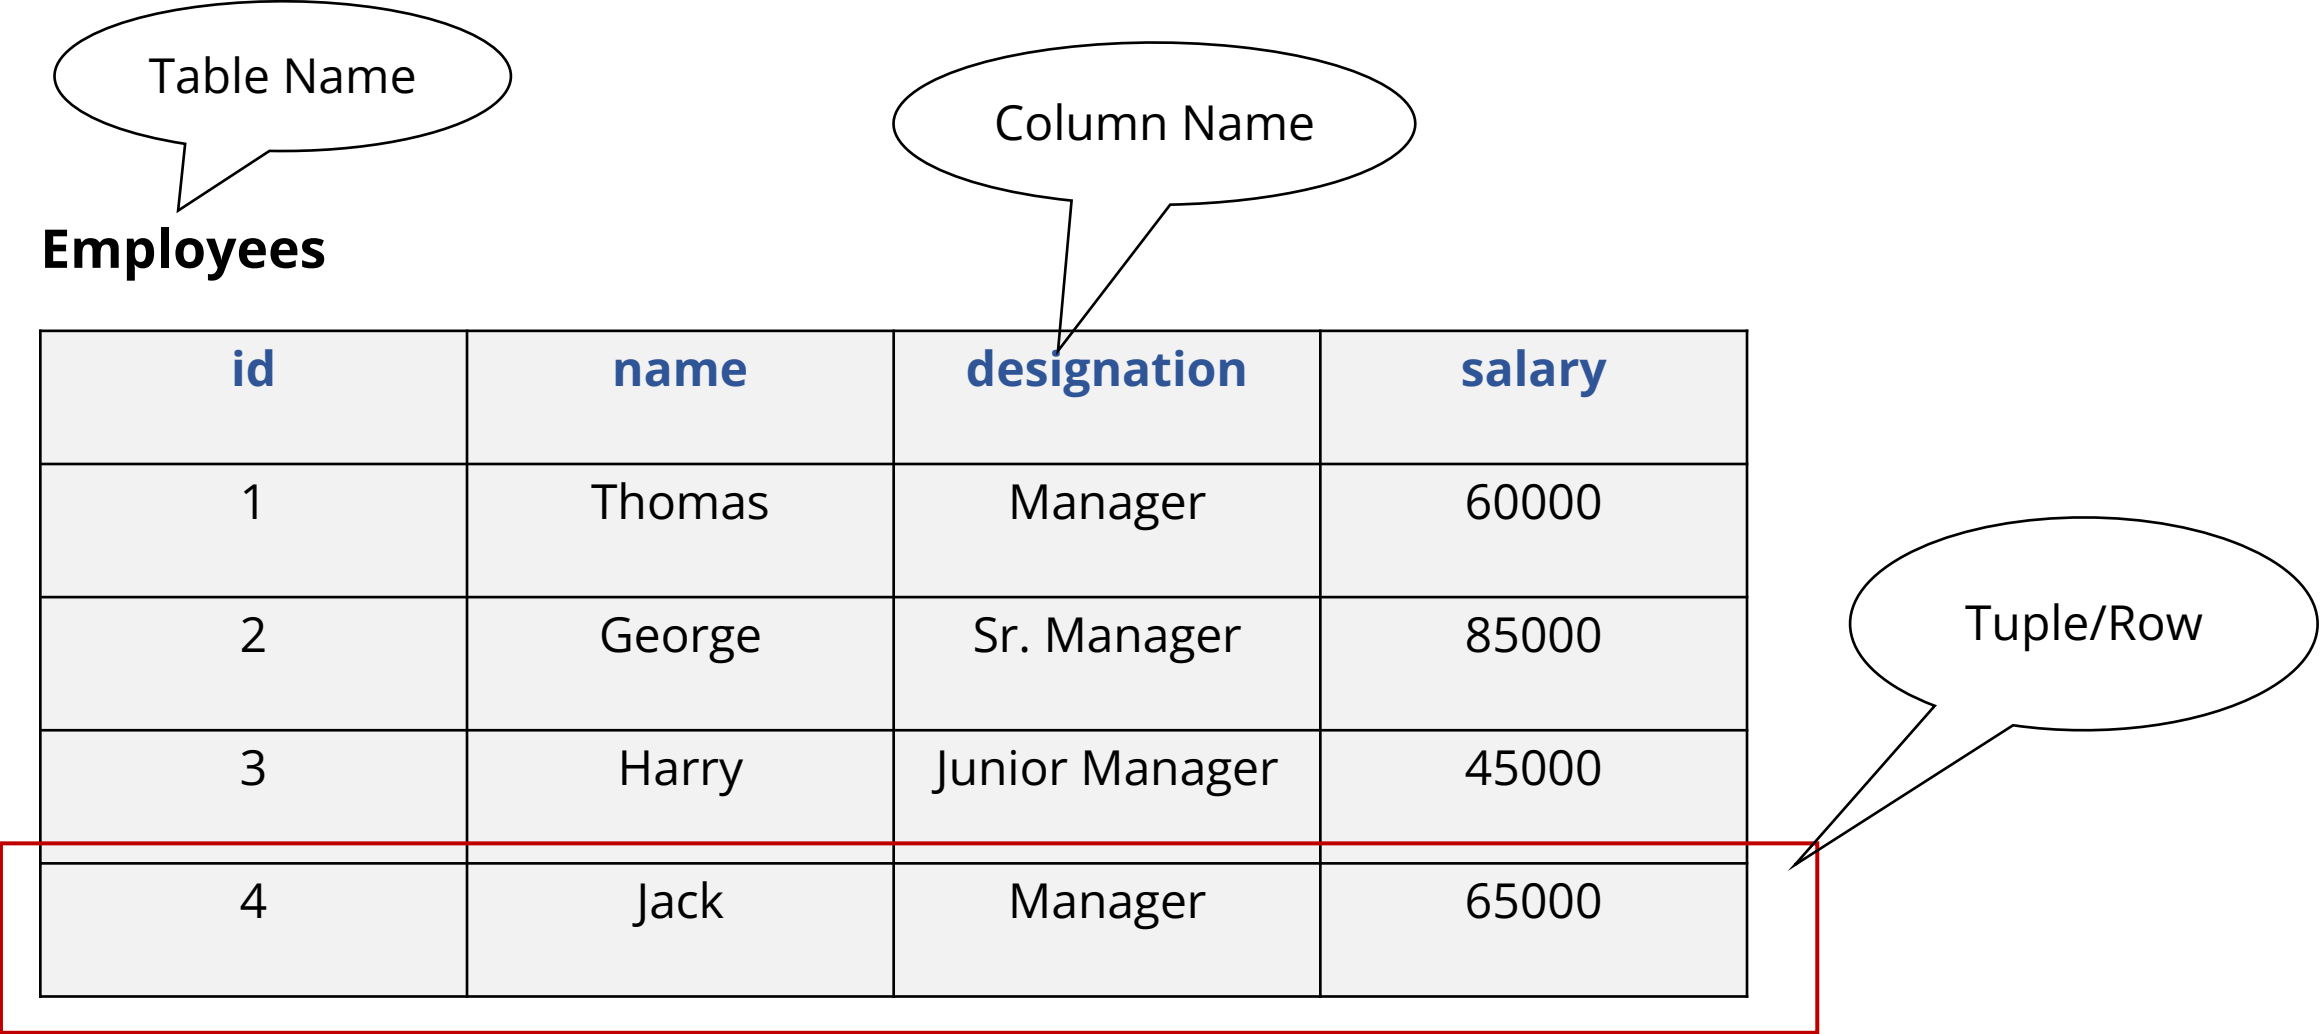
Table Name

Column Name

**Employees**

id	name	designation	salary
1	Thomas	Manager	60000
2	George	Sr. Manager	85000
3	Harry	Junior Manager	45000
4	Jack	Manager	65000

Tuple/Row



# Data Types in SQL

- String data types
  - CHAR, VARCHAR, TEXT
- Numeric data types
  - INT, FLOAT, DOUBLE, BIGINT
- Other data types
  - DATE, DATETIME, TIMESTAMP, YEAR

# SQL Query Form

1.**SELECT** [**DISTINCT**] Attribute\_List **FROM** R1,R2....RM

2.[**WHERE** condition]

3.[**GROUP BY** (Attributes)[**HAVING** condition]]

4.[**ORDER BY**(Attributes)[**DESC**]];



# Sample SQL Queries

**Employees**

id	name	designation	salary
1	Thomas	Manager	60000
2	George	Sr. Manager	85000
3	Harry	Junior Manager	45000
4	Jack	Manager	65000



```
SELECT name, salary  
FROM employees
```

name	salary
Thomas	60000
George	85000
Harry	45000
Jack	65000

# Sample SQL Queries

**Employees**

id	name	designation	salary
1	Thomas	Manager	60000
2	George	Sr. Manager	85000
3	Harry	Junior Manager	45000
4	Jack	Manager	65000



```
SELECT name, salary  
FROM employees  
WHERE salary > 50000
```

name	salary
Thomas	60000
George	85000
Jack	65000

# Sample SQL Queries

**Employees**

id	name	designation	salary
1	Thomas	Manager	60000
2	George	Sr. Manager	85000
3	Harry	Junior Manager	45000
4	Jack	Manager	65000



```
SELECT designation, AVG(salary)
FROM employees
WHERE salary > 50000
GROUP BY designation
ORDER BY designation DESC
```

designation	AVG(salary)
Sr. Manager	85000
Manager	62500

# Spatial SQL

- Uses the same elements and structure of normal SQL
- Allows to work with geospatial types such as geometries and geographies

## Why Spatial SQL?

- Accessible to wider community
- Versatility with many supporting databases and data warehouses
- Multi-dimensional spatial indexing and built-in functions for managing geometry operations
- Efficiency in everyday workflows and task management
- Cross functionality in the organization
- Work with large scale data in SQL enabled data warehouses

# Why Spatial SQL?

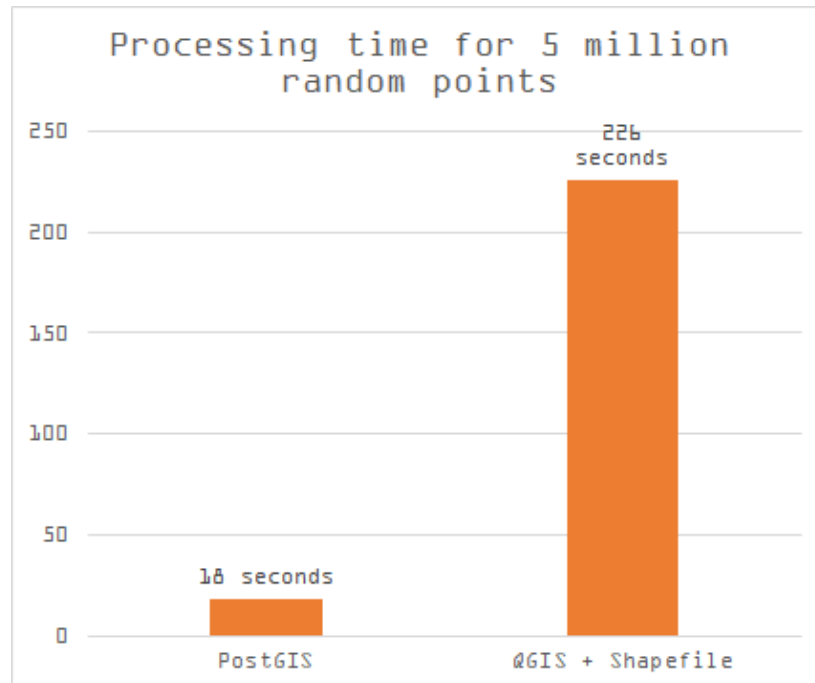
## Efficiency in everyday workflows and task management

- No need to load the same data every time you start a project
- Create new features, join and aggregate data on the fly
- Create indices on geometries to run geometry operations faster
- Update the same table as new data is available
- Create your own user defined functions

# Why Spatial SQL?

## Work with large scale data in SQL enabled data warehouses

- Can speed up processing significantly
- An experiment of creating points in PostGIS take less than 10% of the time taken by QGIS



# Supports for Spatial SQL

## Databases

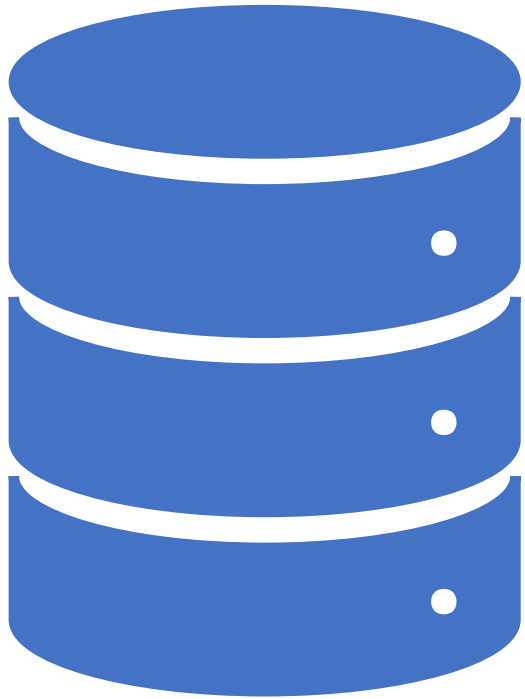
- PostgreSQL with PostGIS
- Microsoft SQL Server
- MySQL
- SQLite with Spatialite
- Oracle Spatial

## Data Warehouses

- Google BigQuery
- Snowflake
- AWS Redshift

## Spatial Extensions of SparkSQL

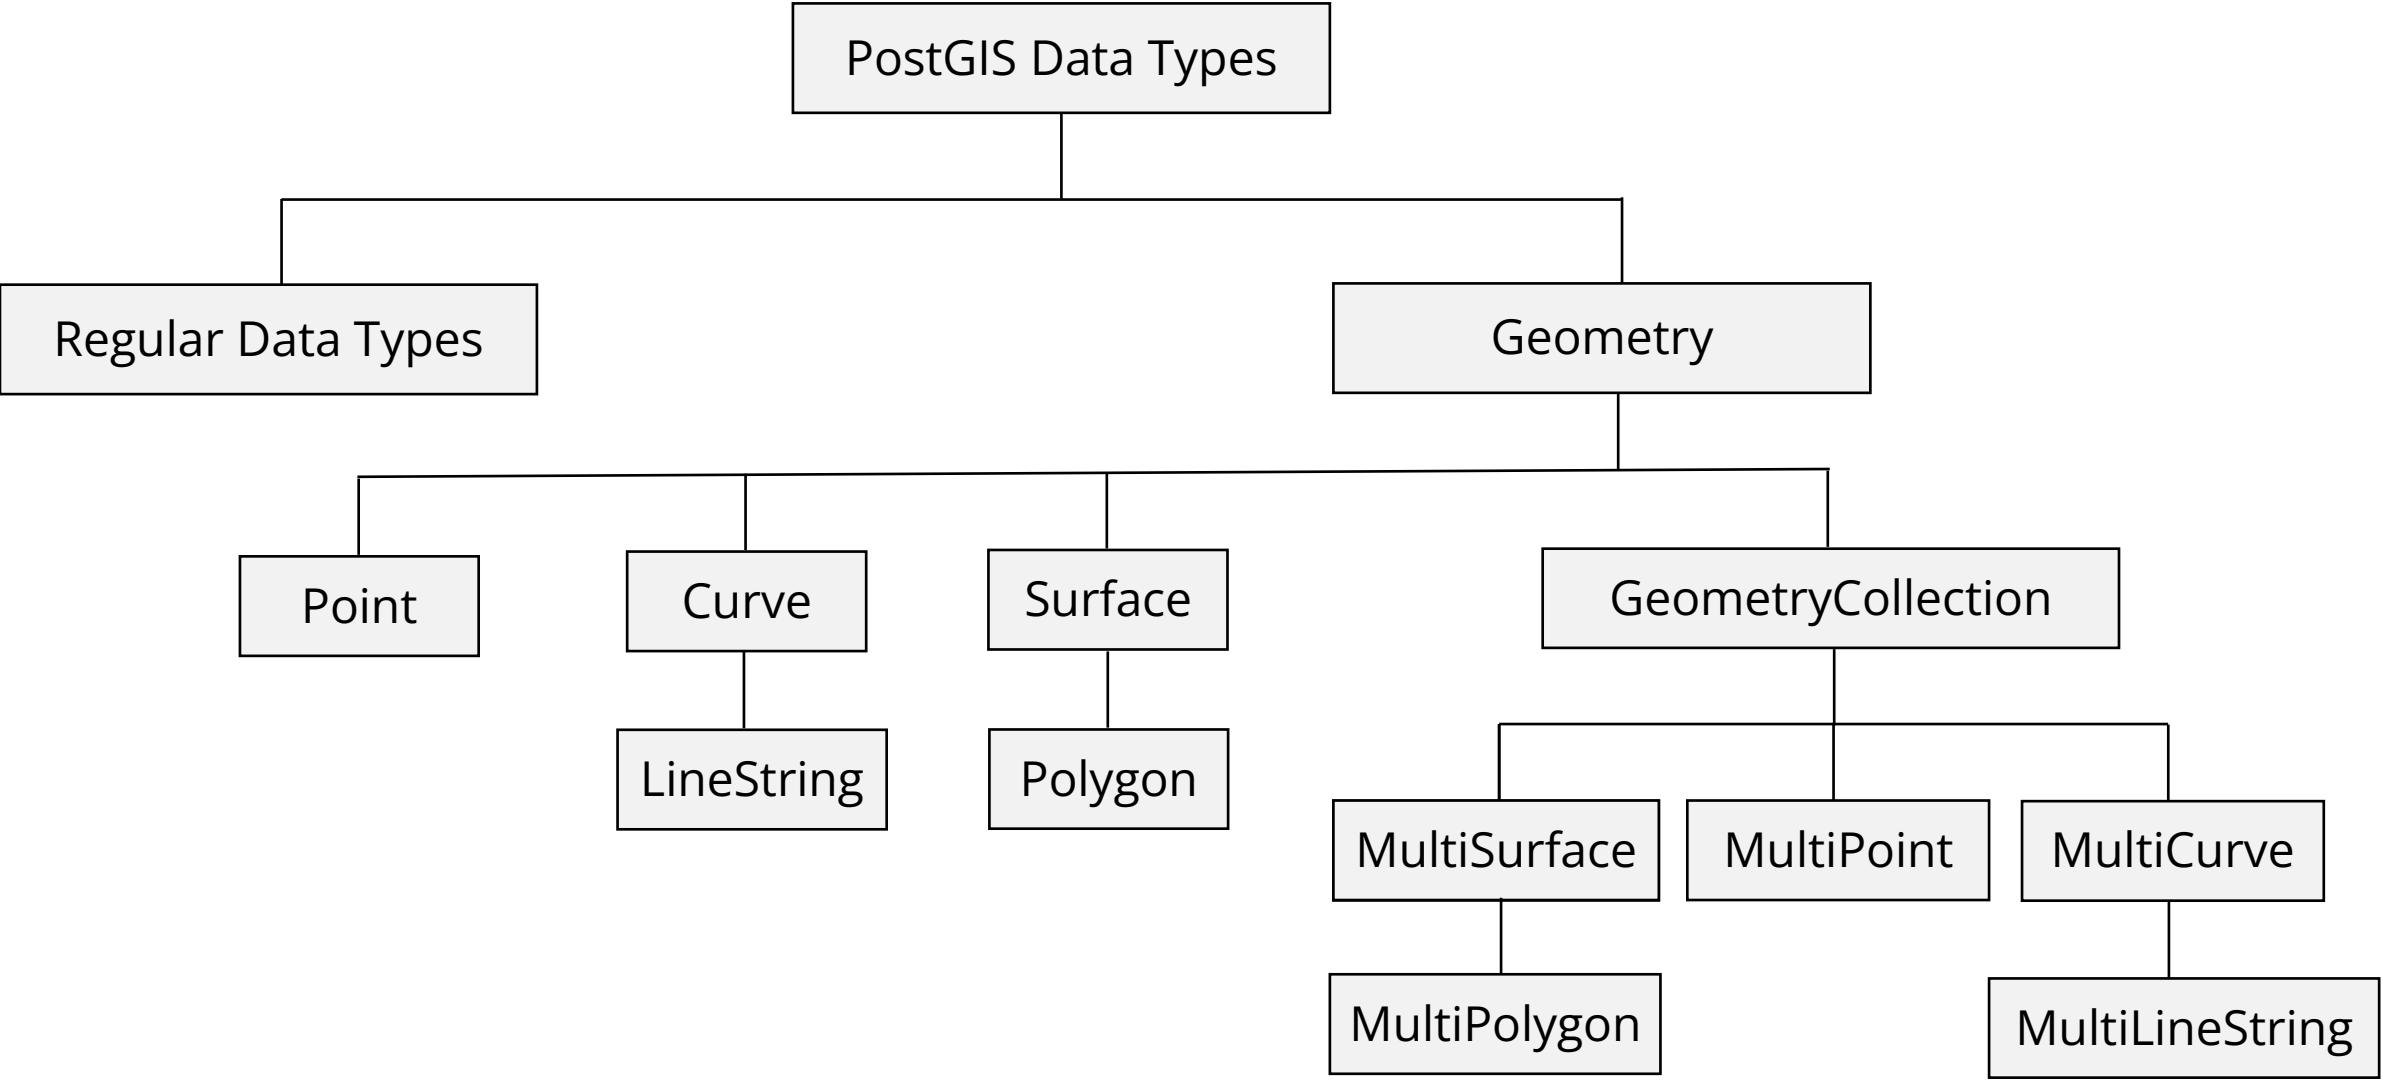
- Apache Sedona, formerly known as GeoSpark
- GeoMesa



# PostGIS Data Model



# PostGIS Data Types



# PostGIS Representation of Geometry Data Type

## Point

- A 0-dimensional geometry that represents a single location in coordinate space

```
POINT (1 2)
```

## LineString

- A 1-dimensional line formed by contiguous sequences of line segments
- Each line segment is defined by two points, with the end point of one segment forming the starting point of second segment

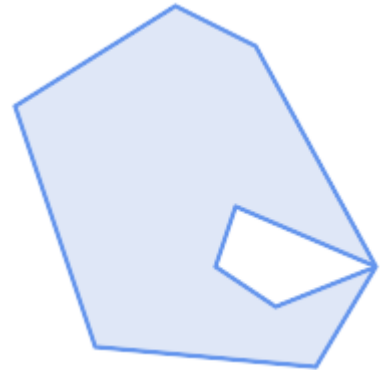
```
LINESTRING (1 2, 3 4, 5 6)
```

# PostGIS Representation of Geometry Data Type

## Polygon

- A 2-dimensional planar region, delimited by an exterior boundary and zero or more interior boundaries (holes)

```
POLYGON ((0 0 0,4 0 0,4 4 0,0 4 0,0 0 0),(1 1 0,2 1 0,2 2 0,1 2 0,1 1 0))
```



## MultiPoint

- A collection of Points

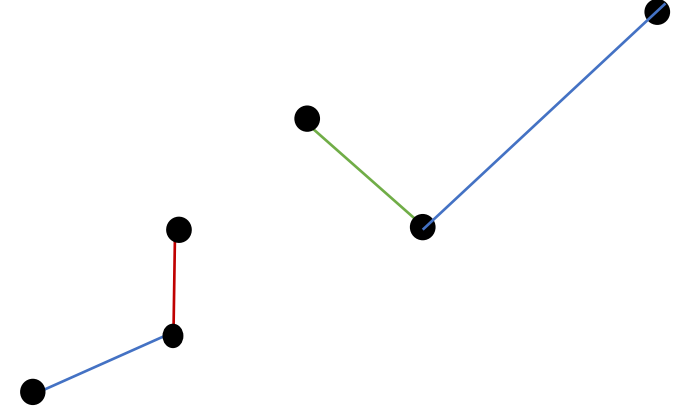
```
MULTIPOINT ( (0 0), (1 2) )
```

# PostGIS Representation of Geometry Data Type

## MultiLineString

- A collection of LineStrings

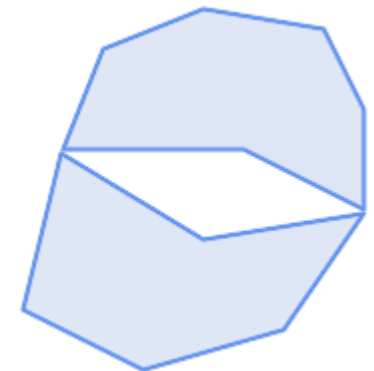
```
MULTILINESTRING ( (0 0,1 1,1 2), (2 3,3 2,5 4) )
```



## MultiPolygon

- A collection of non-overlapping, non-adjacent polygons

```
MULTIPOLYGON (((1 5, 5 5, 5 1, 1 1, 1 5)), ((6 5, 9 1, 6 1, 6 5)))
```



# PostGIS Representation of Geometry Data Type

## GeometryCollection

- A heterogeneous or mixed collection of geometries

```
GEOMETRYCOLLECTION ( POINT(2 3), LINESTRING(2 3, 3 4))
```

## Triangle

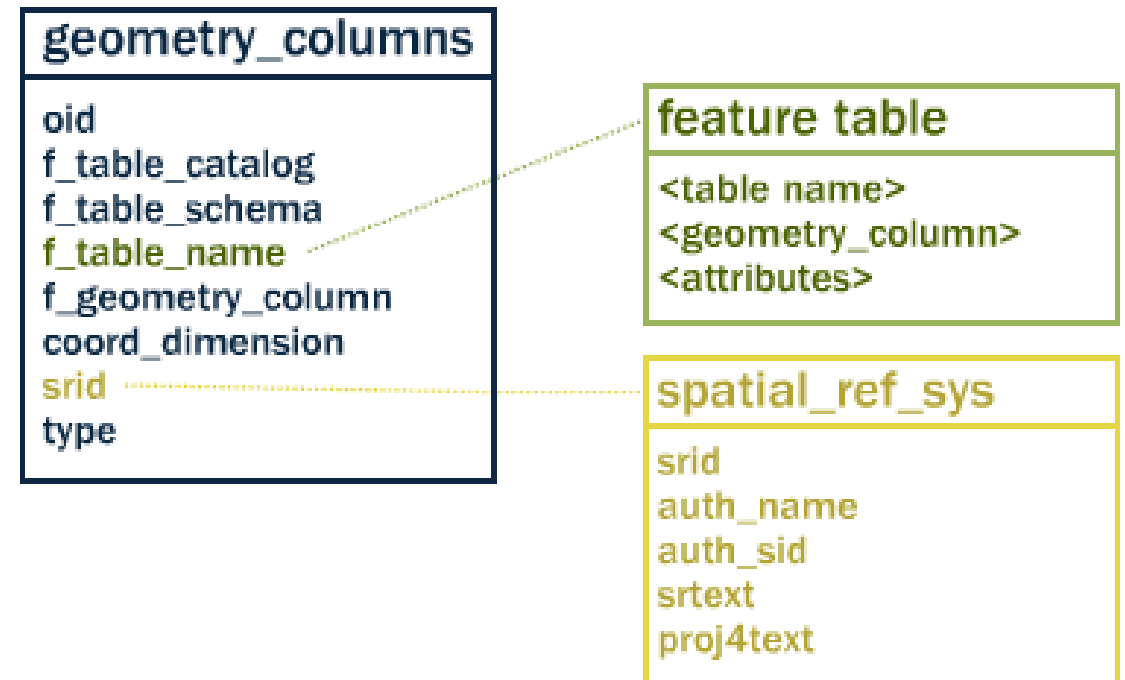
- A Polygon defined by three distinct, non-collinear vertices
- Specified by four coordinates with first and fourth being equal

```
TRIANGLE ((0 0, 0 9, 9 0, 0 0))
```

# PostGIS Metadata Tables

- Table `spatial_ref_sys` defines all spatial reference systems known to the database
- Table `geometry_columns` provides a listing of all features and details of those features
- `f_table_catalog`, `f_table_schema`, `f_table_name` provide the full name of a feature table containing a geometry
- `f_geometry_column` is the name of the geometry column
- `coord_dimension` and `srid` define the dimension of the geometry and spatial reference system identifier
- The `type` column defines the type of geometry

## Table Relationships



# PostGIS SQL

## Creating Table with Geometry Type Column

```
CREATE TABLE TABLE_NAME (name varchar, GEOMETRY_COLUMN_NAME geometry);
```

 ← Default srid is 0

```
CREATE TABLE TABLE_NAME (name varchar, GEOMETRY_COLUMN_NAME geometry(POINT));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOMETRY_COLUMN_NAME geometry(POINT, 4267));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOMETRY_COLUMN_NAME geometry(LINESTRING, 4267));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOMETRY_COLUMN_NAME geometry(POLYGON, 4267));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOMETRY_COLUMN_NAME geometry(POINTZ, 3005));
```

# PostGIS SQL

## Input Output Conversions for WKB and WKT Spatial Objects

```
byte WKB = ST_AsBinary(geometry);  
  
text WKT = ST_AsText(geometry);  
  
geometry = ST_GeomFromWKB(byte WKB, SRID);  
  
geometry = ST_GeometryFromText(text WKT, SRID);
```

## Insert Geometry Data into Table

```
INSERT INTO TABLE_NAME (name, GEOMETRY_COLUMN_NAME )  
VALUES ('Location-1', ST_GeomFromText('POINT(-126.4 45.32)', 4267));
```



# PostGIS SQL

## Creating Table with Geography Type Column

Default srid is 4326



```
CREATE TABLE TABLE_NAME (name varchar, GEOGRAPHY_COLUMN_NAME geography(POINT));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOGRAPHY_COLUMN_NAME geography(POINT, 4267));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOGRAPHY_COLUMN_NAME geography(LINESTRING, 4267));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOGRAPHY_COLUMN_NAME geography(POLYGON, 4267));
```

```
CREATE TABLE TABLE_NAME (name varchar, GEOGRAPHY_COLUMN_NAME geography(POINTZ, 3005));
```

# PostGIS SQL

## Insert Geography Data into Table

```
INSERT INTO TABLE_NAME (name, GEOGRAPHY_COLUMN_NAME )  
VALUES ('Location-1', 'SRID=4326;POINT(-126.4 45.32)');
```

# PostGIS SQL Functions

## Related to Metadata

- ST\_GeometryType(geometry) :- Returns the type of a geometry
- ST\_NDims(geometry) :- Returns the number of dimensions in a geometry
- ST\_SRID(geometry) :- Returns the spatial reference identifier number of a geometry

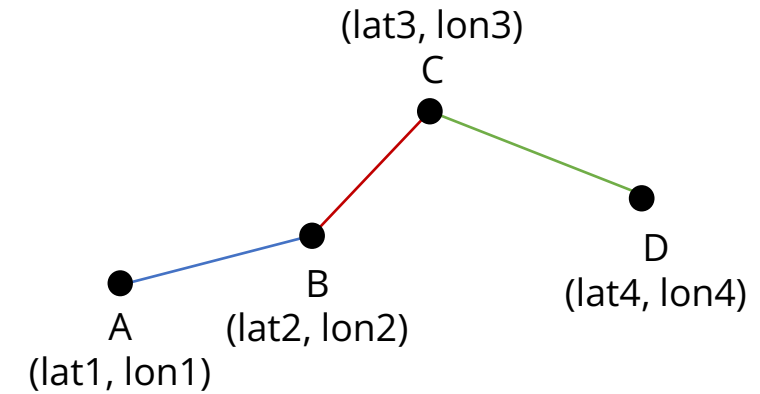
## Related to Point Coordinates

- ST\_X(geometry) :- Returns the X-coordinate of a point geometry
- ST\_Y(geometry) :- Returns the Y-coordinate of a point geometry

# PostGIS SQL Functions

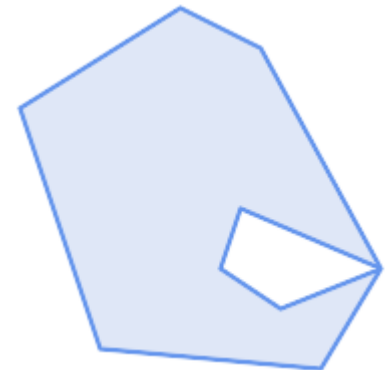
## Related to LineString

- `ST_Length(geometry)` :- Returns the length of the LineString
- `ST_StartPoint(geometry)` :- Returns the first coordinate as a Point
- `ST_EndPoint(geometry)` :- Returns the last coordinate as a Point
- `ST_NPoints(geometry)` :- Returns the number of coordinates in the LineString



## Related to Polygon

- `ST_Area(geometry)` :- Returns the area of the polygon
- `ST_NRings(geometry)` :- Returns the number of rings (1 if there are no holes)
- `ST_ExteriorRing(geometry)` :- Returns the outer ring as a linestring



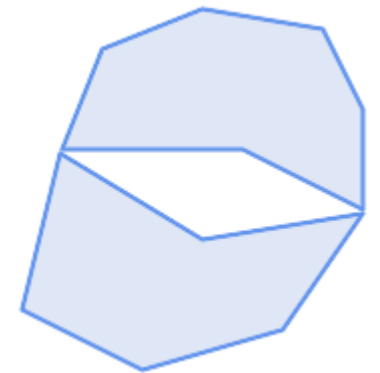
# PostGIS SQL Functions

## Related to Polygon (Continued...)

- `ST_InteriorRingN(geometry, n)` :- Returns a specified interior ring as a linestring
- `ST_Perimeter(geometry)` :- Returns the length of all rings

## Related to Collections

- `ST_NumGeometries(geometry)` :- Returns the number of parts in the collection
- `ST_GeometryN(geometry, n)` :- Returns the specified part
- `ST_Area(geometry)` :- Returns the total area of all Polygonal parts
- `ST_Length(geometry)` :- Returns the total length of all linear parts



# Creating PostGIS Geometry Objects

## Creating Points

```
geometry ST_MakePoint(float x, float y);  
geometry ST_MakePoint(float x, float y, float z);
```

## Creating Lines

```
geometry ST_MakeLine(geometry geom1, geometry geom2);  
geometry ST_MakeLine(geometry[] geoms_array);  
geometry ST_MakeLine(geometry set geoms);
```

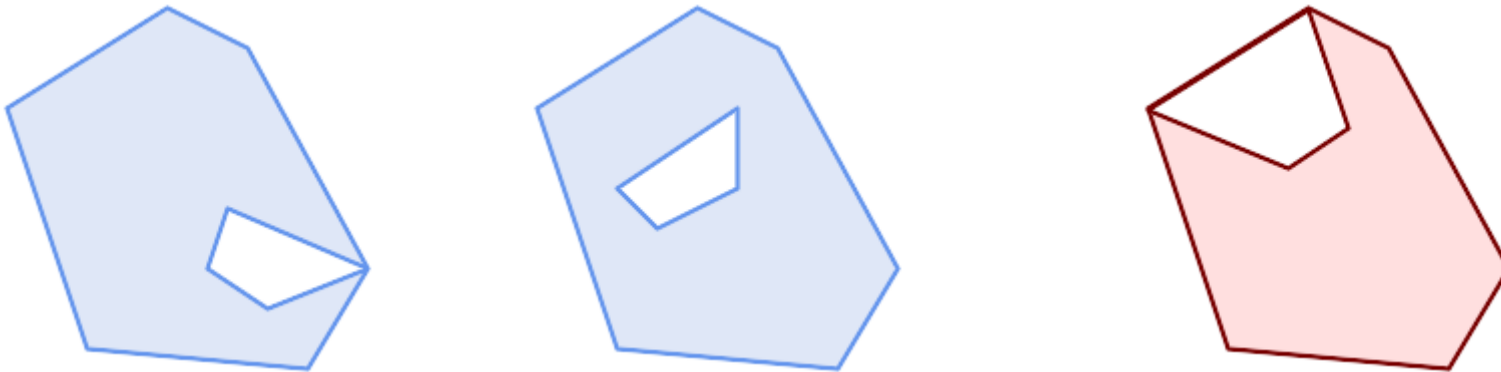
## Creating Polygons

```
geometry ST_MakePolygon(geometry linestring);  
geometry ST_MakePolygon(geometry outerlinestring, geometry[] interiorlinestrings);
```

# Validity of Geometries

## Properties of a Valid Polygon

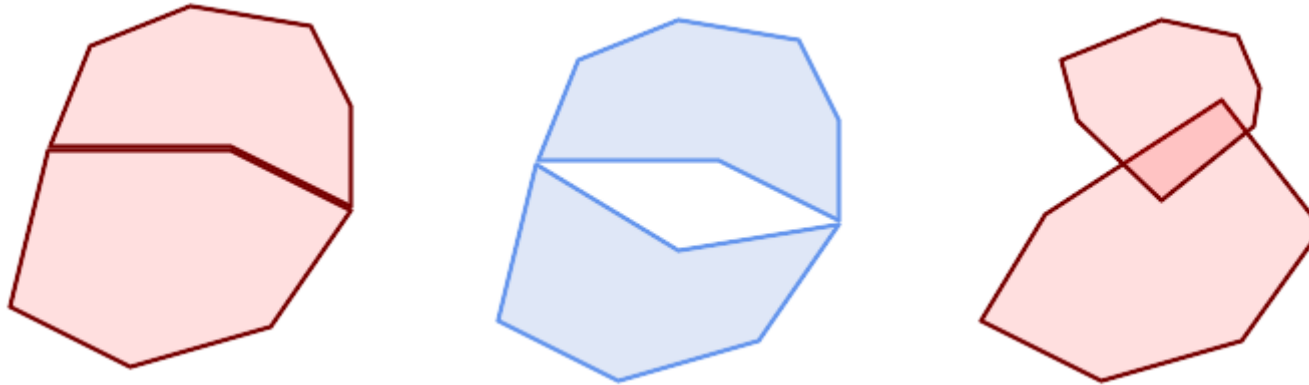
- The boundary rings do not cross and self touch
- The boundary rings may touch at points only as a tangent, not in a line
- Interior rings are within the exterior ring
- The polygon interiors should not touch in a way that splits the polygon into parts



# Validity of Geometries

## Properties of a Valid MultiPolygon

- Element polygons are valid
- Elements do not overlap or their interiors must not intersect
- Elements touch only at points, not along a line



Source: [https://postgis.net/docs/using\\_postgis\\_dbmanagement.html#LinearRing](https://postgis.net/docs/using_postgis_dbmanagement.html#LinearRing)





# All ST Functions Supported by PostGIS

Visit the link:

<https://postgis.net/docs/manual-1.5/ch08.html>



# Apache Sedona Spatial SQL Data Model



# Supported Geometry Objects

- Point
- MultiPoint
- LineString
- MultiLineString
- Polygon
- MultiPolygon

# Sedona SQL Data Structure

- Datasets are represented as Spark DataFrames
- Each DataFrame can be considered as a Table in PostGIS
- Each attribute is a column, while each data instance is a row
- Spatial DataFrame contains a geometry type column

# Spatial Operation Through SQL with Spatial DataFrames

## A Sample Spatial DataFrame: dfSpatialSample

Name	Area	Geometry
Name-1	470	POLYGON ((933100.92 192536.09, 933091.01 192572.17, 933088.58 192604.97, 933779.28 195908.73, 933841.76 195957.79, 933100.92 192536.09))
Name-2	520	MULTIPOLYGON (((1033269.24 172126.00, 1033439.64 170883.95, 1033473.26 170808.21, 1033269.24 172126.00)), ((1033422.35 157944.65, 1033419.99 157936.99, 1033408.21 157938.17, 1033422.35 157944.65)))
Name-3	300	POLYGON ((933100.92 192536.09, 933091.01 192572.17, 933088.58 192604.97, 933779.28 195908.73, 933841.76 195957.79, 933100.92 192536.09))
Name-4	740	POLYGON ((933100.92 192536.09, 933091.01 192572.17, 933088.58 192604.97, 933779.28 195908.73, 933841.76 195957.79, 933100.92 192536.09))

# Spatial Operation Through SQL with Spatial DataFrames

## Running SQL Queries on dfSpatialSample

1. Create a temporary view from the DataFrame object

```
dfSpatialSample.createOrReplaceTempView("sample_spatial_view")
```

2. Run SQL queries assuming the view name as a table name

```
dfSpatialSample = sparkSession.sql("SELECT ... FROM sample_spatial_view WHERE ...")
```

3. Display the schema

```
dfSpatialSample.printSchema()
```

# Constructors for Creating Geometry Objects

- ST\_geomFromText
- ST\_GeomFromWkB
- ST\_Point
- ST\_GeomFromGeoJSON
- ST\_PolygonFromEnvelop

## Sample Usage

```
SELECT ST_GeomFromWKT('POINT(40.7128 -74.0060)') AS geometry
```

More details: <https://sedona.apache.org/api/sql/Constructor/>

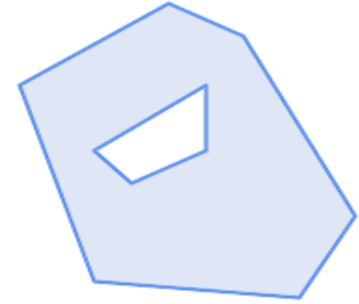
# Predicates Supported by Sedona SQL

- ST\_Within(A, B) – returns True if A is fully contained by B
- ST\_Disjoint(A, B) – returns True if A and B are disjoint
- ST\_Intersects(A, B) – returns True if A intersects B

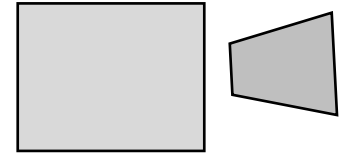
## Sample Usage

```
SELECT * FROM pointdf  
WHERE ST_Intersects(ST_PolygonFromEnvelope(1.0,100.0,1000.0,1100.0),  
                    pointdf.arealandmark)
```

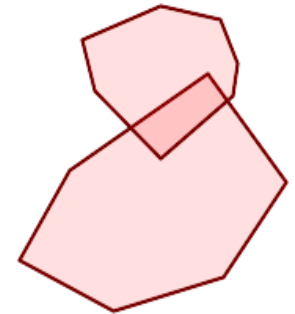
More details: <https://sedona.apache.org/api/sql/Predicate/>



**Within**



**Disjoint**



**Intersects**



# Sedona SQL Functions

- ST\_Distance(A, B) – returns the Euclidean distance between A and B
- ST\_StartPoint(A) – returns the first point of a given linestring
- ST\_GeometryType(A) – returns the type of the geometry as a string

## Sample Usage

```
SELECT ST_GeometryType(polygondf.countyshape) FROM polygondf
```

More details: <https://sedona.apache.org/api/sql/Function/>