

CSE 594: Spatial Data Science & Engineering

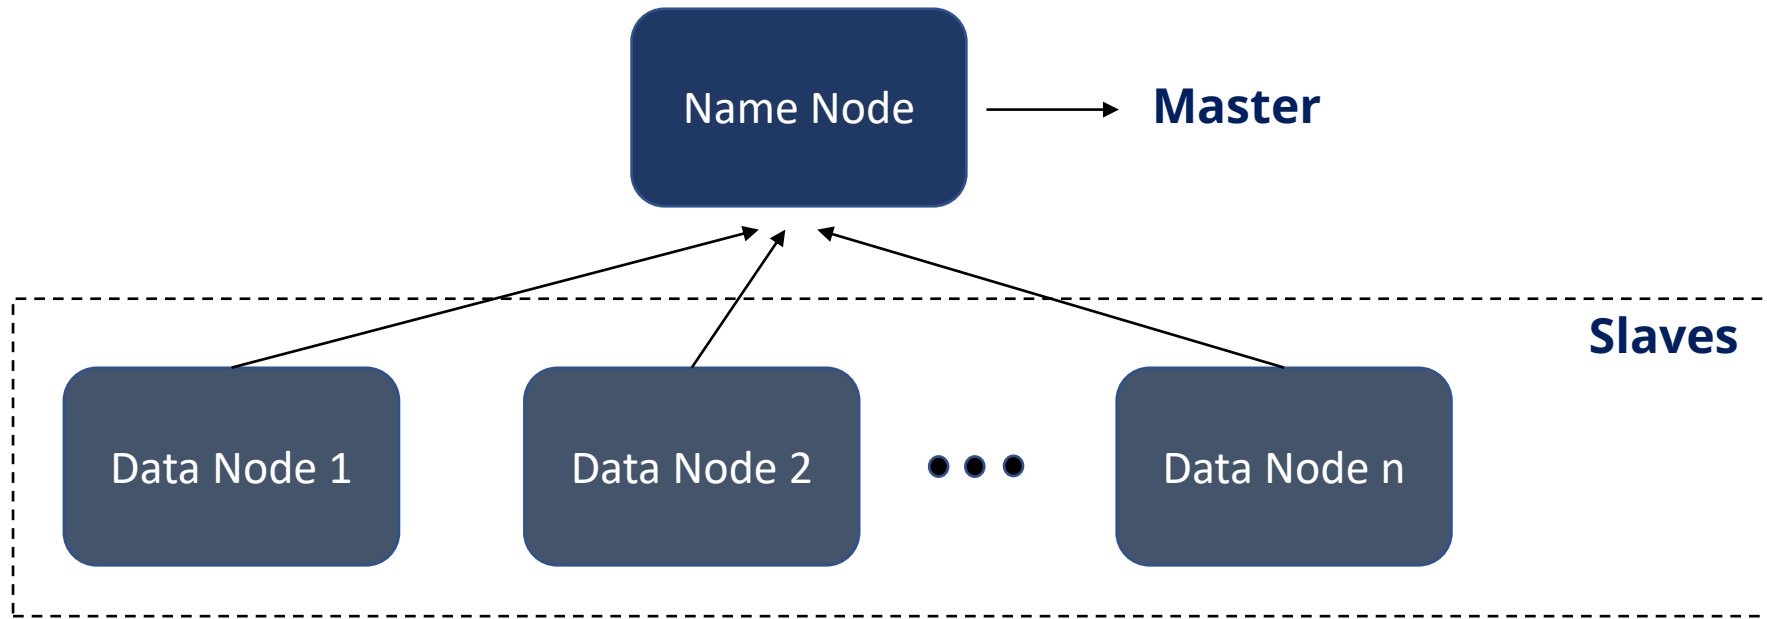
Lecture 9

Big Spatial Data Systems

Apache Hadoop

- A framework that uses distributed storage and parallel processing to store and manage big data
- An ecosystem consisting of three components:
 - ❑ Hadoop HDFS – the storage unit
 - ❑ Hadoop YARN – resource management unit
 - ❑ Hadoop MapReduce – data processing unit

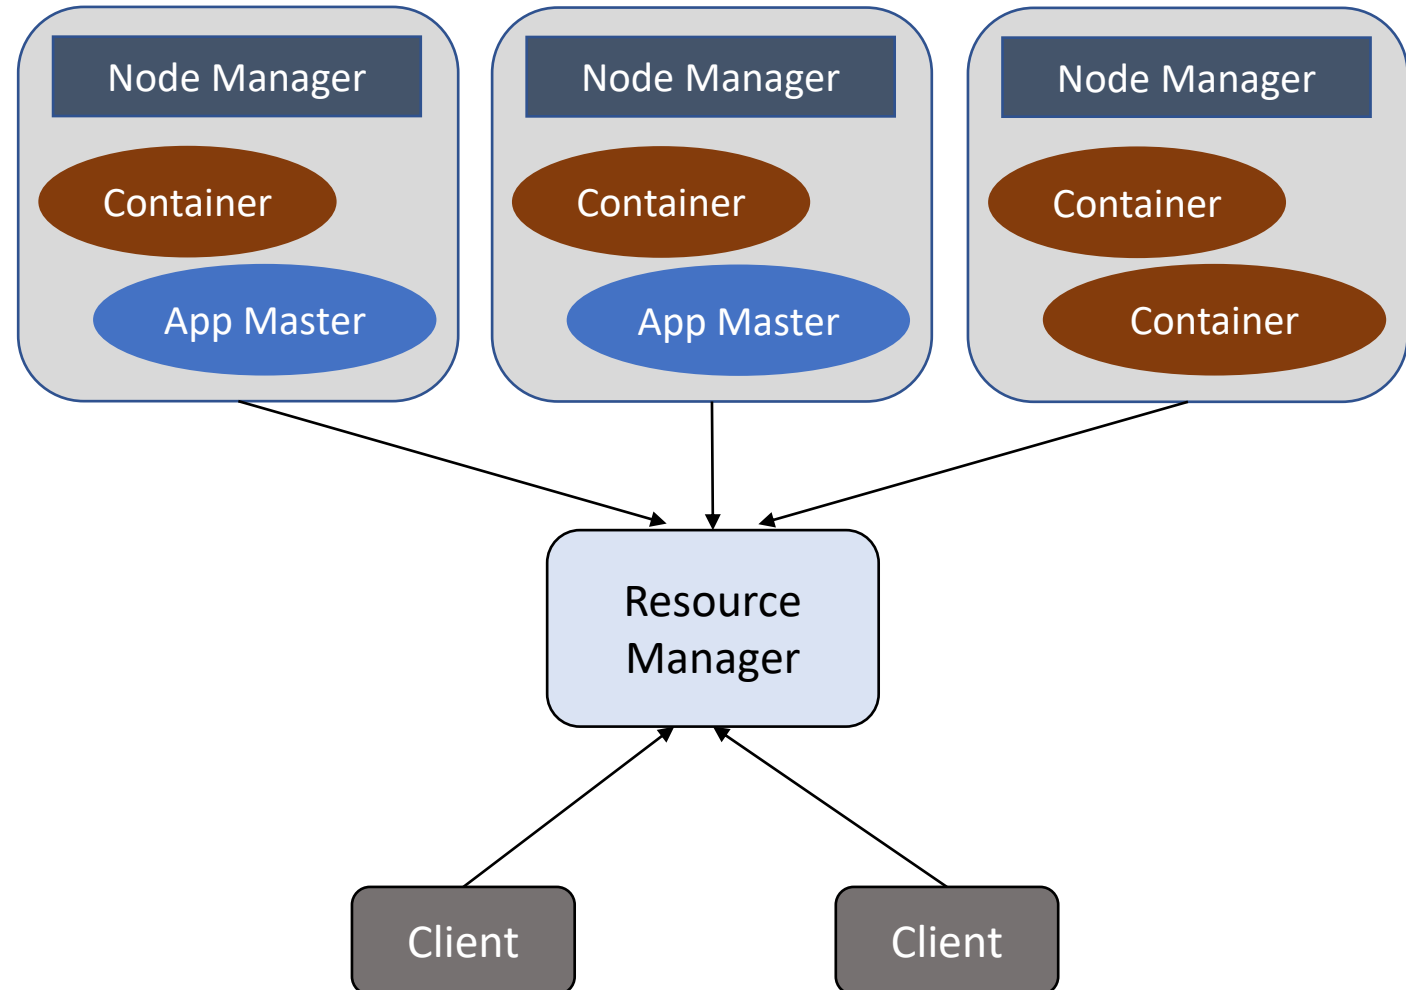
HDFS – Hadoop Distributed File System



- Two components – one name node and multiple data nodes
- Data is divided into 128 MB blocks and distributed across data nodes
- Name node stores the meta data and manages data nodes
- Data nodes read, write, process and replicate the data, and send signals to the name node

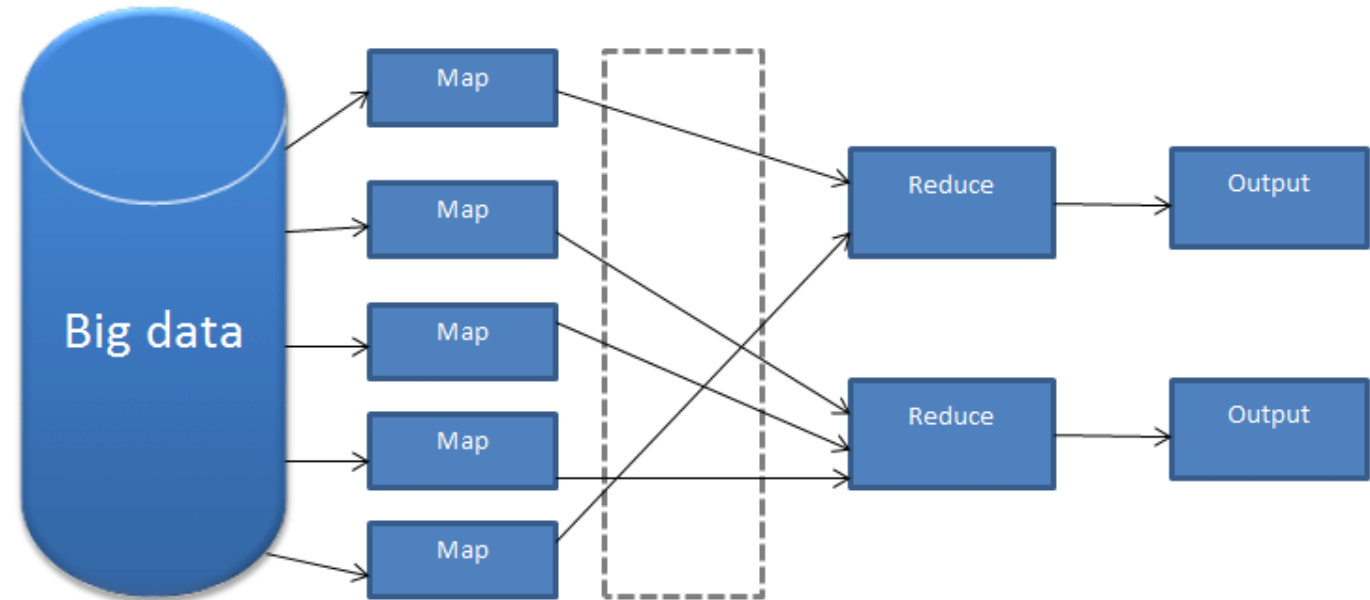
Hadoop YARN – Yet Another Resource Negotiator

- Acts like an OS to Hadoop.
- Another file system on top of HDFS
- Responsible for managing cluster resources
- Performs job scheduling
- Allows Hadoop to run different distributed applications simultaneously
- Separates the processing layer from resource management layer

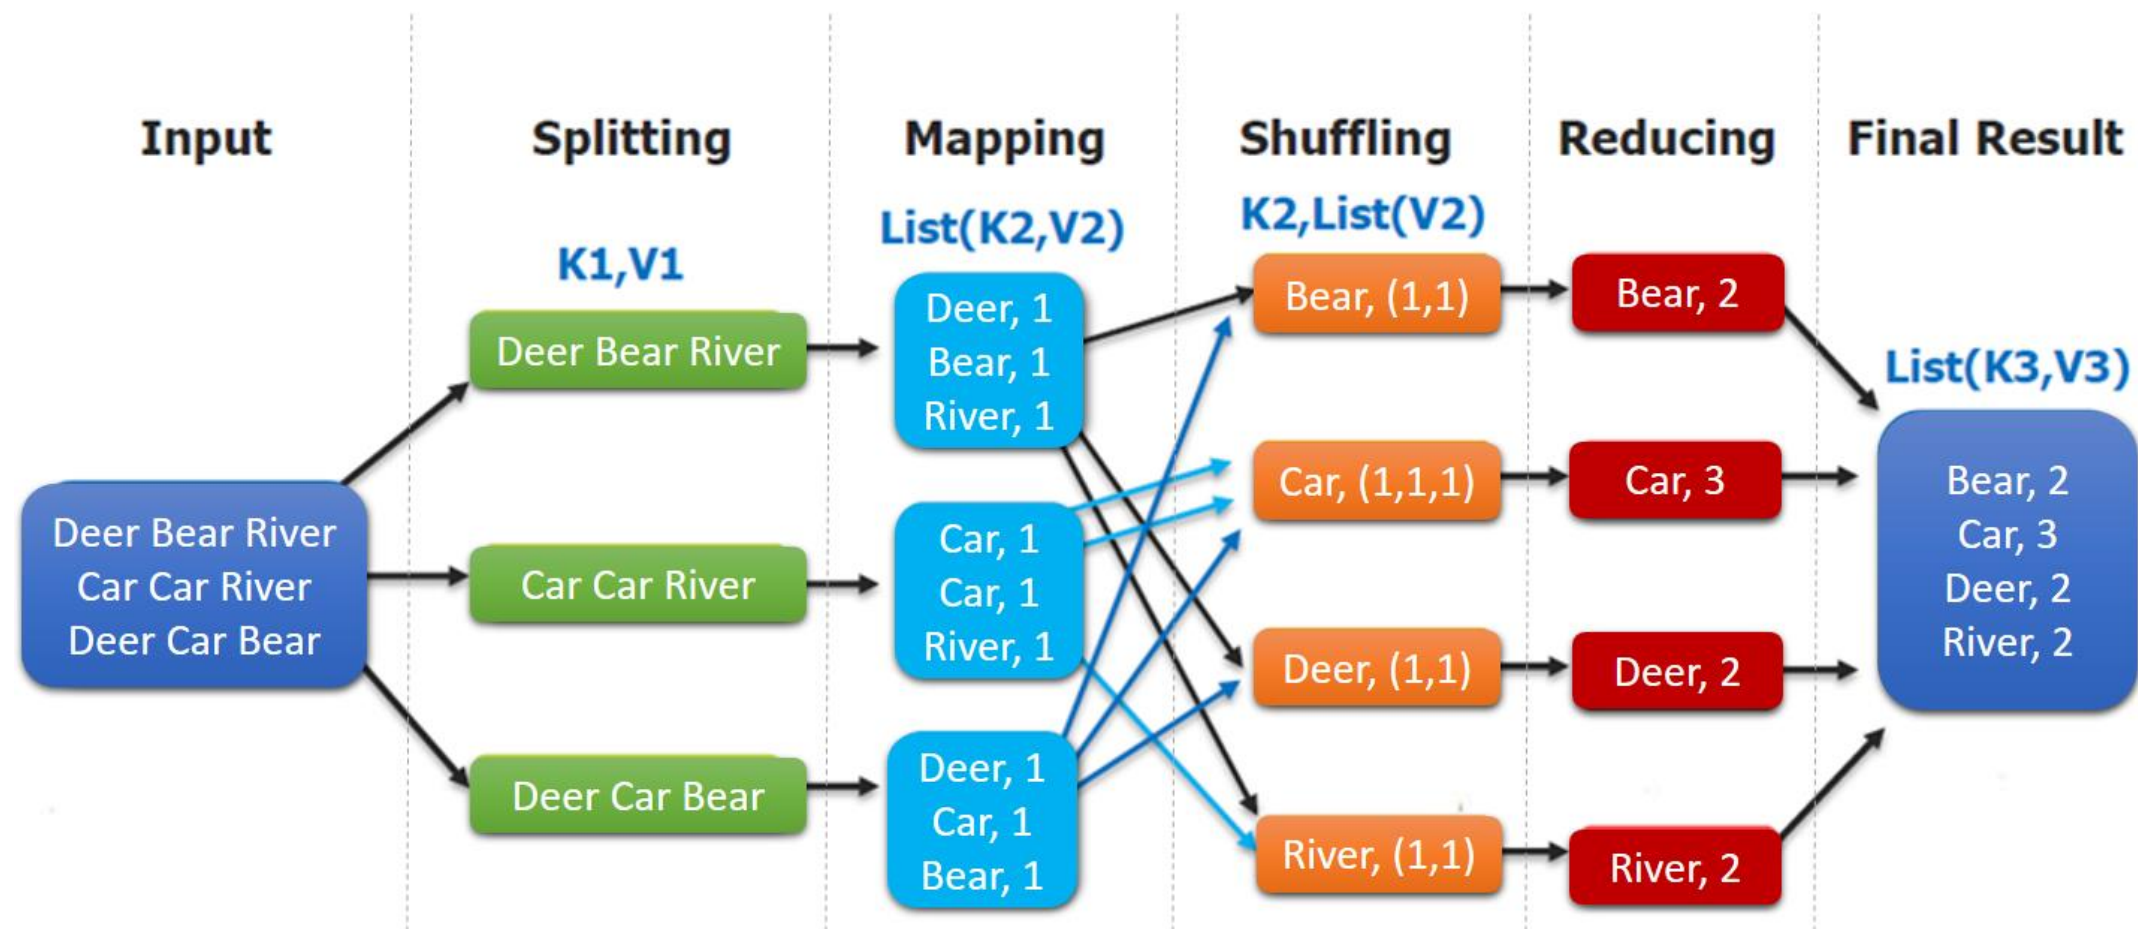


Hadoop MapReduce

- A programming model for parallelizing tasks into multiple nodes
- Consists of two phases – mapping phase followed by reducing phase
- Mapper reads and processes the data to produce key-value pairs
- Reducer takes the outputs from mappers and aggregates the key-value pairs to produce the final output

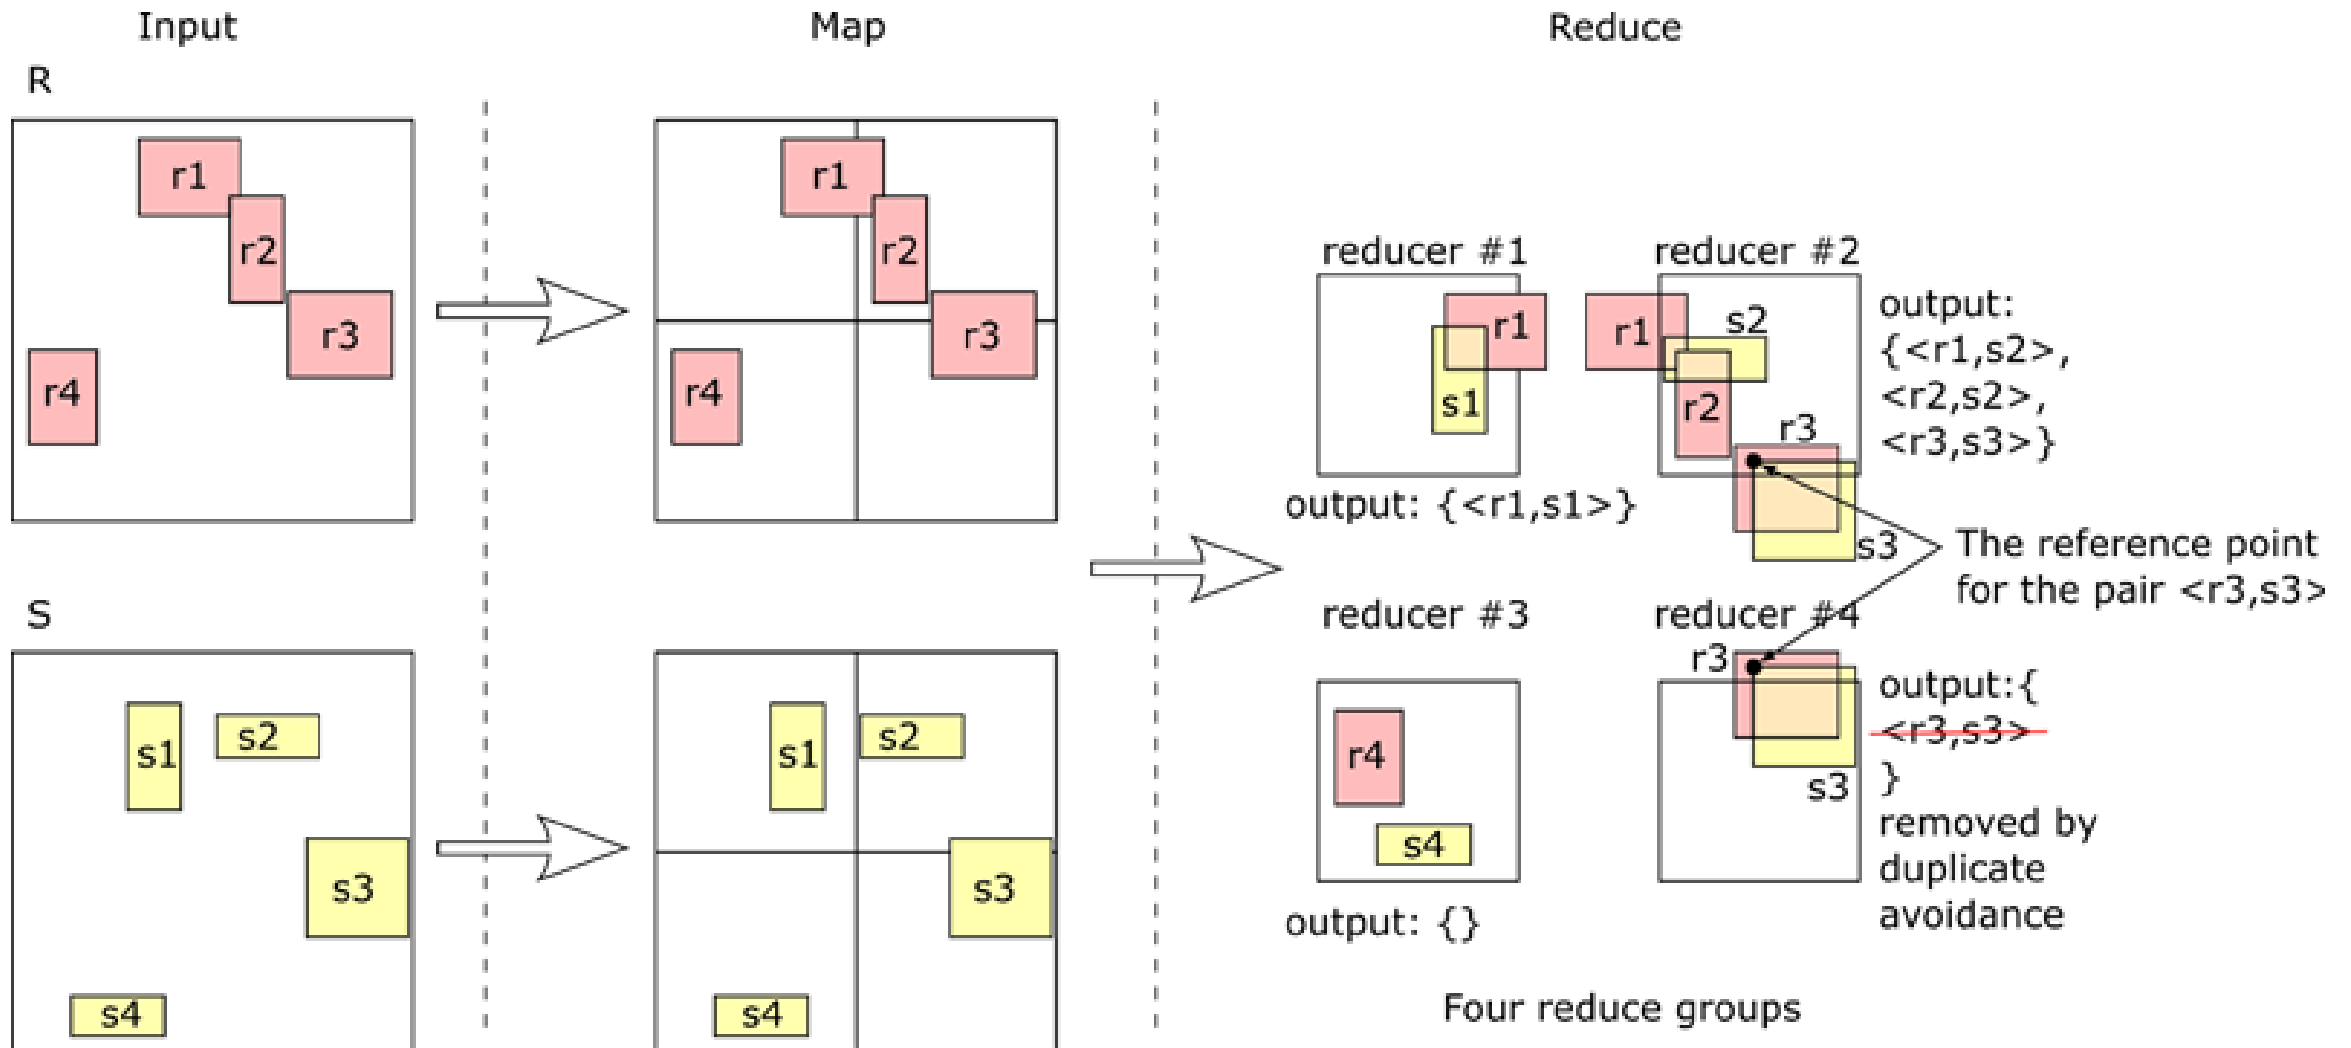


Simple MapReduce Example



MapReduce Algorithm for Word Counting Problem

MapReduce for Spatial Join



MapReduce for Spatial Join

- Partitions the input rectangles according to a fixed sized grid
- Store the records in the corresponding cells
- Records overlapped with multiple cells need to be replicated in all overlapping cells
- Perform the join in each machine in parallel
- A mapper reports the pairs of cell ids and spatial objects in the corresponding partition
- A reducer collects the pairs from mappers and performs the join operation followed by a duplicate removal step

Removing Duplicates

- Compute the reference point (top left corner of the intersection) for a candidate joining pair
- Report a pair only if the reference point lies within the grid cell

MapReduce for Spatial Operations

- Some spatial operations may need only one of either mapper or reducer in practice

Computing Dataset Boundary

- Needs only reduce operation
- Calculate the merged rectangular boundary of spatial objects two by two until all objects are aggregated
- After calculating the boundary for each partition, aggregate the boundaries of partitions

Transforming Reference System

- Needs only mapping operation
- For each partition, traverse the objects in the corresponding partition and convert their SRS

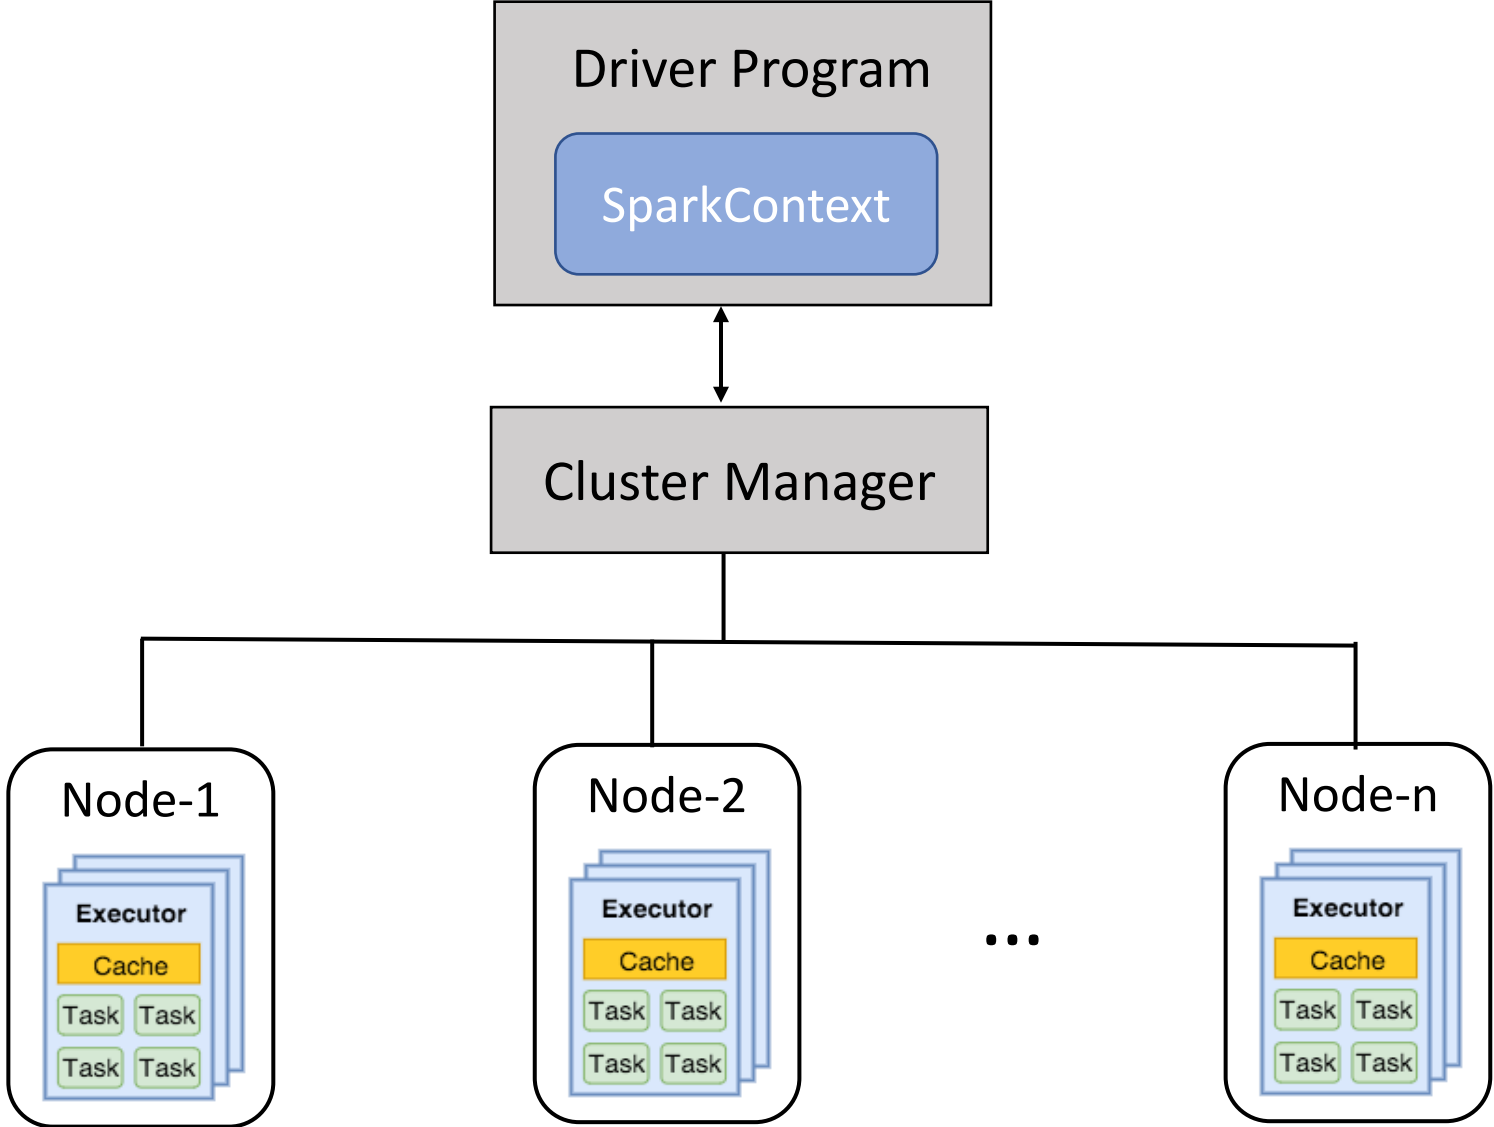
Apache Spark

- Distributed Data Processing Framework
- Can perform processing tasks on very large datasets
- Distribute data processing tasks across multiple computers
- Supports in-memory data caching and reuse across computations
- Supports majority of the data formats and various storage systems

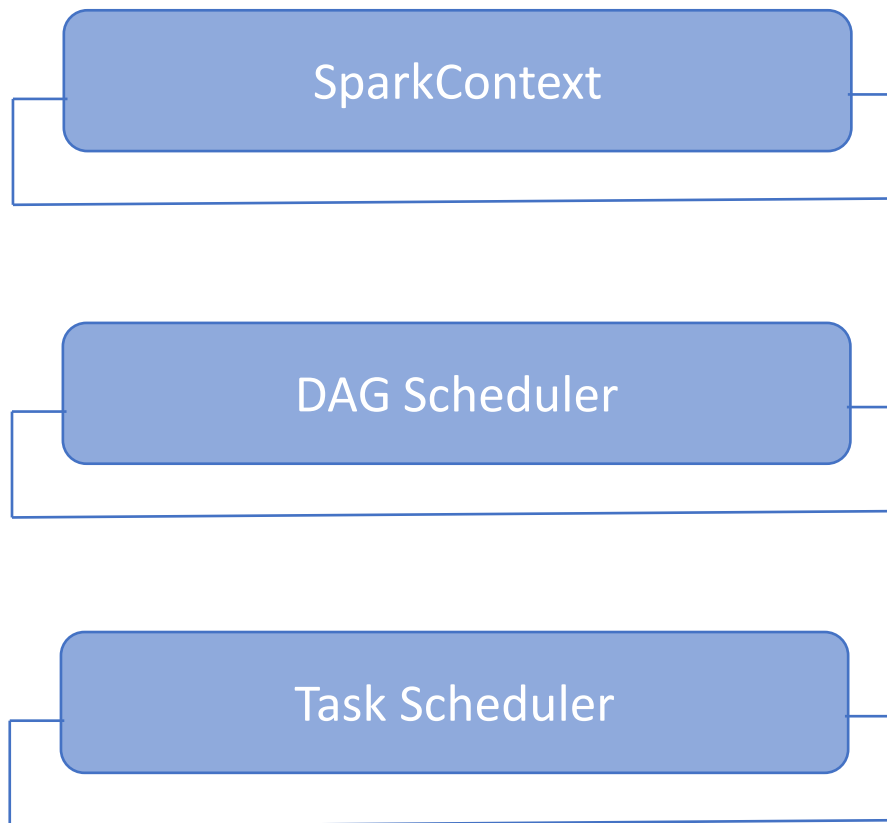
Spark Key Concepts

- Job
- Stages
- Tasks
- DAG
- Executor

Spark Components

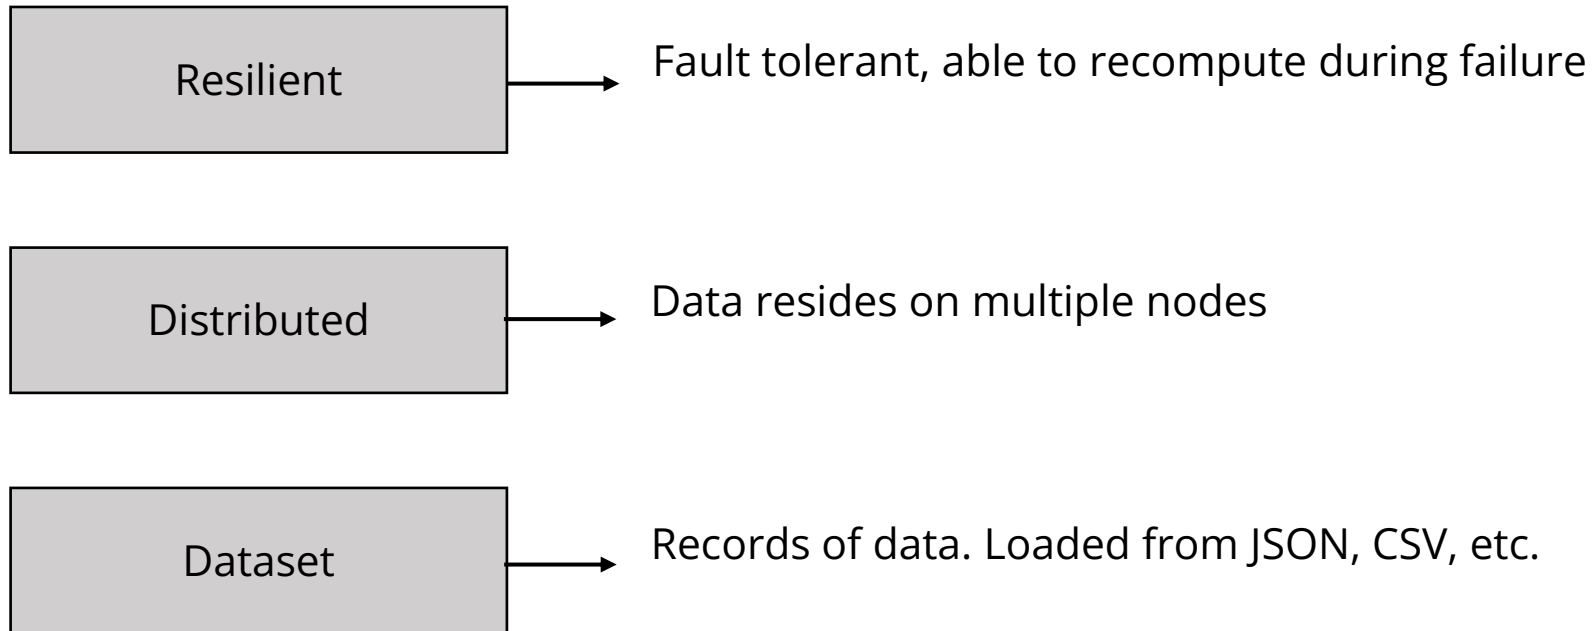


Spark Driver Components



Spark RDD

- Resilient Distributed Dataset
- Spark representation of a set of data



An RDD is an immutable distributed collection of objects. Each RDD is split into multiple partitions and can be computed on different nodes

Properties of Spark RDD

- In-memory computation
- Lazy evaluation
- Fault tolerance
- Immutable
- Partitioning
- Persistence
- Coarse-grained operations

Spark RDD Operations

- **Transformations**

Takes an RDD as input and produce one or more RDDs as output

- **Actions**

RDD operations that produce non-RDD values. One of the ways to send results from executors to the driver.

- Collect()
- Count()

Advantages of Apache Spark

- Speed
 - Spark's in memory data engine can perform tasks up to hundred times faster than Hadoop map-reduce
 - Very efficient for tasks that require frequent write operations
 - Even if data cannot completely be contained in memory it is 10 times faster than map-reduce
- Developer friendly API
- Advanced DAG execution engine
- Numerous storage engines: HDFS, Cassandra, HBase, etc.

Challenges of Handling Spatial Objects with Spark RDDs

- Heterogeneous data sources
 - ❑ Spatial data is stored in a variety of formats – CSV, GeoJSON, WKT, shape files, GeoTIFF, GeoParquet
- Complex geometrical shapes
 - ❑ Different objects in a dataset may have a variety of shapes – concave/convex polygon, multi-polygon, line, point
- Spatial partitioning
 - ❑ Default partitioner in Spark does not preserve spatial proximity
- Spatial index support
 - ❑ Spatial indexes are different from regular indexes used by Spark

Apache Sedona

- A cluster computing system for processing large-scale spatial data, initially released as GeoSpark
- Add spatial support to Spark by solving the challenges – heterogeneous data sources, complex geometry shapes, spatial partitioning, and spatial index support
- Extends Apache Spark with a set of Spatial RDDs that efficiently load, process, and analyze large-scale spatial data across machines
- Facilitate to create spatial analytics and data mining applications and run them in any Spark environments

Why Apache Sedona?

- **High Speed**

2X – 10X faster than other Spark-based geospatial data systems on computation intensive query workloads

- **Low Memory Consumption**


50% less peak memory consumption than other Spark based geospatial data systems

- **Ease of Use**

Spatial SQL API has supports for Scala, Java and Python. API is very easy to use

What Makes Apache Sedona Efficient?

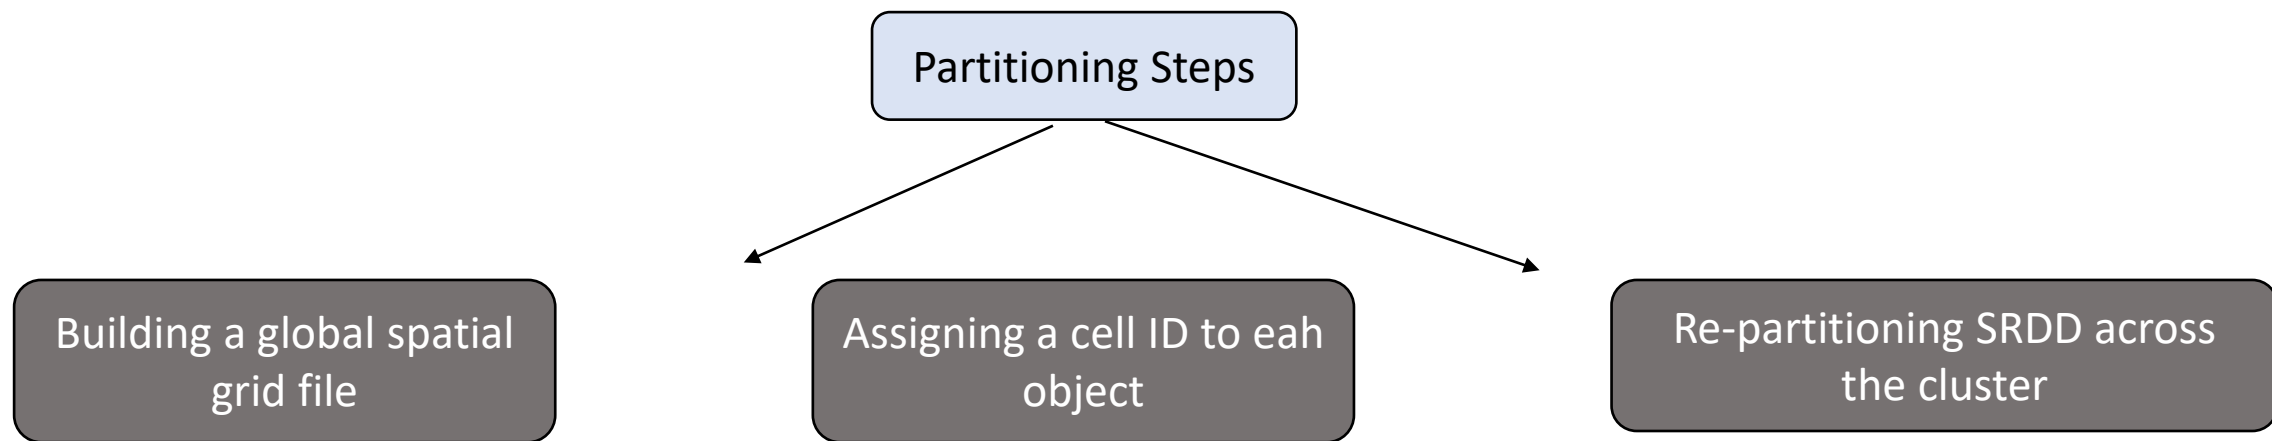
- Introduces Spatial Resilient Distributed Dataset (SRDD) which provides in-built support for geometrical and distance operations
- Spatial RDD can accommodate heterogeneous spatial objects which are very common in a GIS area
- Introduces spatial query processing layer which is optimized for spatial query operators such as range filter, distance filter, spatial k-nearest neighbors, range join and distance join
- When executing a spatial SQL query, the optimizer takes into account the execution time and produces a good query execution plan



Spatial RDD Layer

- Supports loading data from any all popular spatial formats – CSV, WKT/WKB, GeoJSON, Shape file, GeoTIFF, GeoParquet
- Supports heterogeneous geometry shapes – point, multi-point, linestring, multi-linestring, polygon, multi-polygon, circle, and geometry collection
- Provides a built-in library for executing geometrical computation on spatial RDDs
- Provides native support for many common geometrical operations

Spatial RDD Partitioning



Building a Global Spatial Grid File

- Samples a subset of spatial data preserving the spatial distribution
- Create the partitions on the sampled data
- Partitioning on sampled data is used for entire dataset

Spatial RDD Partitioning

Building a Global Spatial Grid File


- Supports grid creation with following approaches

- ☐ Uniform Grid


- ☐ KDB-Tree

- ☐ Quad-Tree


- ☐ R-Tree




(a) SRDD partitioned by uniform grids



(b) SRDD partitioned by Quad-Tree



(c) SRDD partitioned by R-Tree



(d) SRDD partitioned by KDB-Tree

Spatial RDD Partitioning

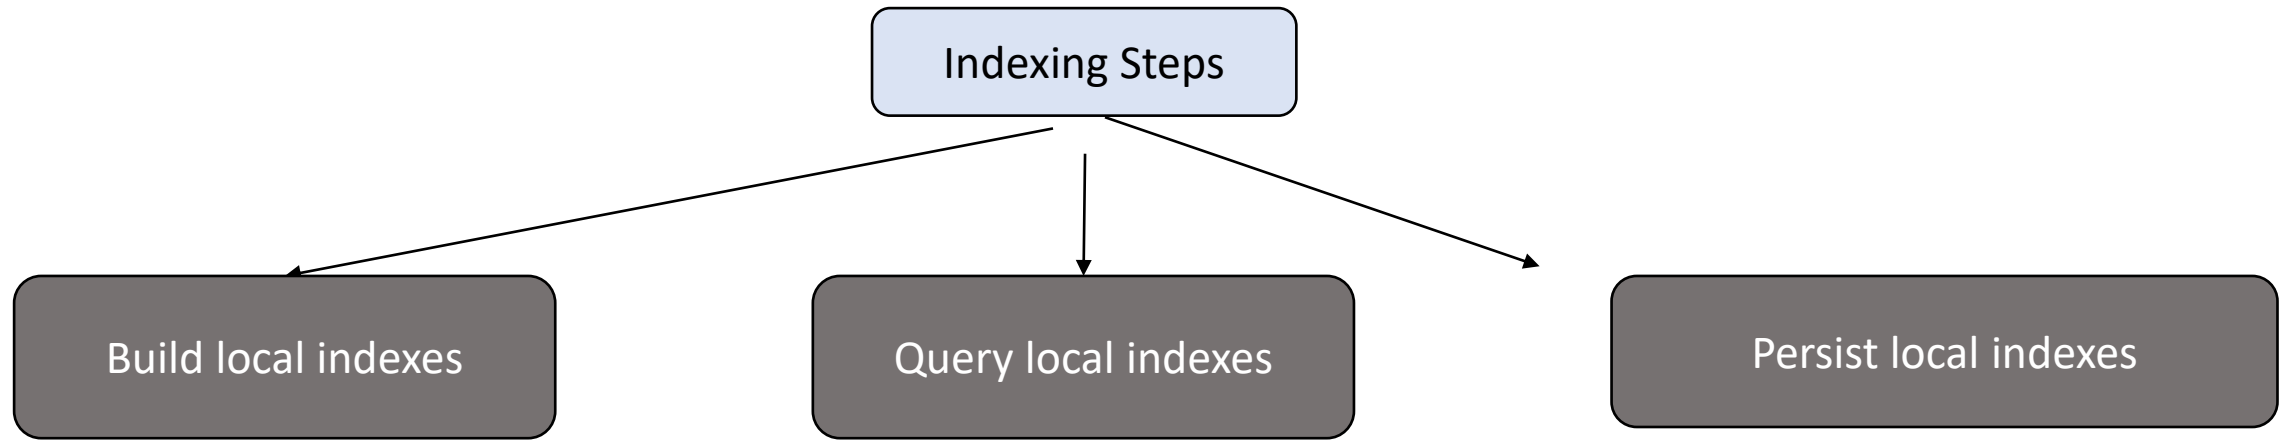
Assigning a Cell ID to Each Object

- Duplicates the grid files and broadcasts the copies to each spatial RDD partition
- After receiving broadcasted grid file, original spatial RDD partitions check each internal object against the grid file
- Results are stored in a new RDD of key-value pairs, where key is the grid ID and object is the spatial object having intersection with the corresponding grid

Re-partitioning SRDD Across the Cluster

- Spatial objects that have the same grid cell ID are grouped into the same partition
- Needs lots of shuffle across the cluster

Spatial RDD Indexing



Spatial RDD Custom Serializer

- Spark default serializer is efficient for regular data types
- Improves over spatial data types for faster data transfer across the cluster

Spatial Query Processing Layer

Processing Spatial Range and Distance Queries

- Applies filter and refinement strategy
- If spatial index exists, use the query window's MBR to query the spatial index and return candidate results
- If no spatial index, filters spatial objects from all objects in a partition using query window
- Check the spatial relation between query window and candidate objects

Processing K Nearest Neighbor Queries

- 2 phases – selection phase and sorting phase
- Selection phase selects K-nearest objects from each partition using filter and refinement model
- Sort all objects selected by various partitions and return to k

Spatial Query Processing Layer

Processing Spatial Join Queries

- Zip partitions from both dataset
- Perform partition level local join
- Perform aggregation of outputs from local joins
- Remove duplicates using reference point approach
- Also support broadcast join for small datasets

Some Supported Spatial Operations

- Spatial Range Query
- Spatial Join
- Spatial KNN Query
- Spatial Indexing

Some Applications

- **Region heat map:** visualize taxi trip pickup points distribution in an entire city on a map
- **Spatial aggregation:** show the taxi trip pickup count distribution per taxi zone
- **Spatial co-location pattern mining:** Are taxi trip pickup points co-located with area landmarks such as airports, museums, hospitals, etc.?